

```
In [3]: # section 2.1 Processing and Import necessary Libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as msno

# Load the dataset
df = pd.read_csv("/content/The Climate Change Twitter Dataset.csv")

# Display first 5 rows
print("First 5 rows of the dataset:")
display(df.head())

# Shape of the dataset
print(f"Dataset contains {df.shape[0]} rows and {df.shape[1]} columns.")

# Data types and missing values
print("\nData info:")
df.info()

# Summary statistics for numeric and categorical features
print("\nSummary statistics (numeric):")
display(df.describe())

print("\nSummary statistics (categorical):")
display(df.describe(include='object'))

# Check missing values
print("\nMissing values count:")
display(df.isnull().sum())

# Visualize missing values
msno.matrix(df)
plt.title("Missing Value Matrix")
plt.show()
```

First 5 rows of the dataset:

	created_at	id	lng	lat	topic	sentiment	stance	gender	temperature_avg	aggressiveness
0	2006-06-06 16:06:42+00:00	6132	NaN	NaN	Weather Extremes	-0.097180	neutral	female	NaN	aggressive
1	2006-07-23 21:52:30+00:00	13275	-73.949582	40.650104	Weather Extremes	0.575777	neutral	undefined	-1.114768	aggressive
2	2006-08-29 01:52:30+00:00	23160	NaN	NaN	Weather Extremes	0.500479	neutral	male	NaN	aggressive
3	2006-11-07 02:46:52+00:00	57868	NaN	NaN	Weather Extremes	0.032816	neutral	male	NaN	aggressive
4	2006-11-27 14:27:43+00:00	304553	NaN	NaN	Importance of Human Intervantion	-0.090428	neutral	male	NaN	aggressive

Dataset contains 15789411 rows and 10 columns.

Data info:

```
<class 'pandas.core.frame.DataFrame'>
```

RangeIndex: 15789411 entries, 0 to 15789410

Data columns (total 10 columns):

```
#   Column      Dtype
---  -
0   created_at  object
1   id          int64
2   lng         float64
3   lat         float64
4   topic       object
5   sentiment   float64
6   stance      object
7   gender      object
8   temperature_avg float64
9   aggressiveness object
```

dtypes: float64(4), int64(1), object(5)

memory usage: 1.2+ GB

Summary statistics (numeric):

	id	lng	lat	sentiment	temperature_avg
<b>count</b>	1.578941e+07	5.307538e+06	5.307538e+06	1.578941e+07	5.307538e+06
<b>mean</b>	8.459853e+17	-4.639117e+01	3.408025e+01	2.536663e-03	1.245156e+00
<b>std</b>	3.113522e+17	7.523162e+01	2.229430e+01	4.379192e-01	3.799786e+00
<b>min</b>	6.132000e+03	-1.796670e+02	-9.000000e+01	-9.942049e-01	-2.328904e+01
<b>25%</b>	7.354169e+17	-9.536327e+01	3.315067e+01	-3.957429e-01	-1.140978e+00
<b>50%</b>	9.564851e+17	-7.703637e+01	3.995233e+01	-2.328273e-03	1.211522e+00
<b>75%</b>	1.049540e+18	-1.483154e-01	4.550884e+01	4.161248e-01	3.867153e+00
<b>max</b>	1.178912e+18	1.793830e+02	8.500000e+01	9.917458e-01	2.100350e+01

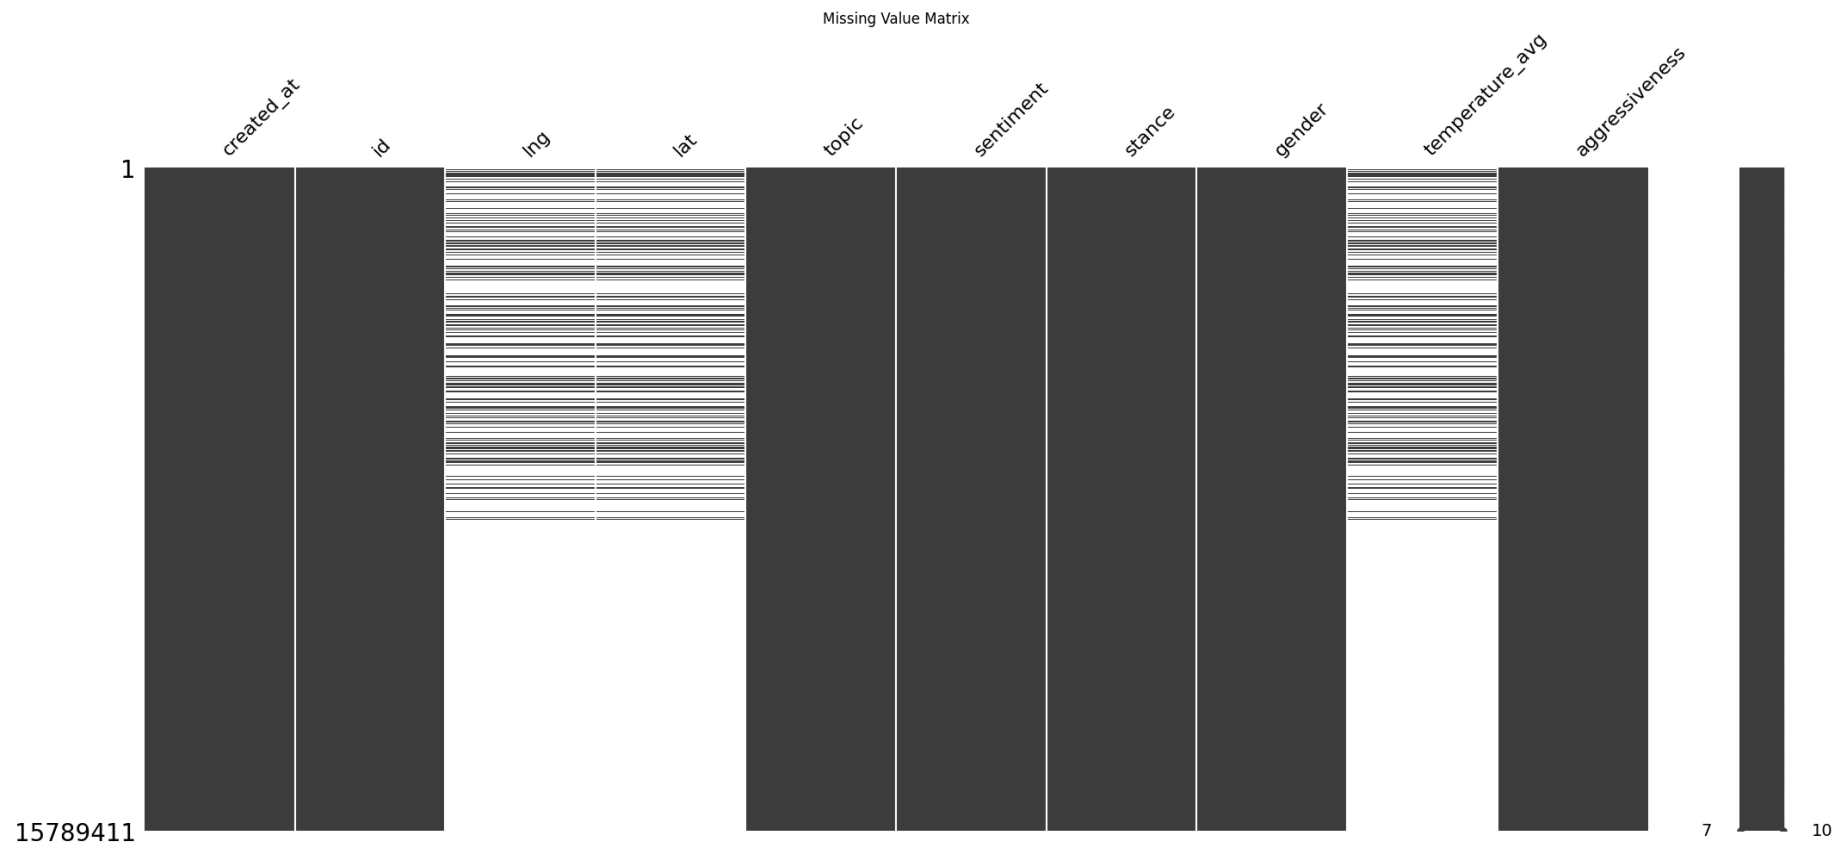
Summary statistics (categorical):

	created_at	topic	stance	gender	aggressiveness
<b>count</b>	15789411	15789411	15789411	15789411	15789411
<b>unique</b>	13390455	10	3	3	2
<b>top</b>	2016-09-25 19:35:06+00:00	Global stance	believer	male	not aggressive
<b>freq</b>	68	4135619	11292424	10307402	11262144

Missing values count:

	0
created_at	0
id	0
lng	10481873
lat	10481873
topic	0
sentiment	0
stance	0
gender	0
temperature_avg	10481873
aggressiveness	0

dtype: int64



```
In [4]: # Drop rows with missing values for key predictive features
df_cleaned = df[['id','created_at','lng','lat','sentiment', 'temperature_avg', 'gender', 'stance', 'topic', 'aggressiveness']]
print("First 5 rows of the dataset:")
df_cleaned['created_at'] = pd.to_datetime(df_cleaned['created_at'], errors='coerce')
display(df_cleaned.head())
```

First 5 rows of the dataset:

	id	created_at	lng	lat	sentiment	temperature_avg	gender	stance	topic	aggressiveness
1	13275	2006-07-23 21:52:30+00:00	-73.949582	40.650104	0.575777	-1.114768	undefined	neutral	Weather Extremes	aggressive
7	1092823	2006-12-14 01:39:10+00:00	-122.419420	37.774930	-0.544195	4.228540	male	neutral	Ideological Positions on Global Warming	aggressive
8	1278023	2006-12-17 19:43:09+00:00	-79.791980	36.072640	-0.565028	5.478175	male	denier	Weather Extremes	aggressive
9	1455543	2006-12-21 01:39:01+00:00	-121.805790	38.004920	0.650960	-1.652156	male	neutral	Weather Extremes	not aggressive
11	1893063	2006-12-31 10:47:25+00:00	-1.902691	52.479699	0.670905	4.864521	male	neutral	Weather Extremes	aggressive

In [ ]:

```
In [5]: #Data Transformation- Encoding Categorical Variables
from sklearn.preprocessing import LabelEncoder

# Initialize dictionary to hold the encoders
encoders = {}

# Iterate over the categorical columns and encode them
for col in ['gender', 'stance', 'topic', 'aggressiveness']:
    le = LabelEncoder()
    original_values = df_cleaned[col].copy() # Save original values for display
    encoded_column_name = f"{col}_encoded" # Name for encoded column

    # Create a new column for the encoded values
    df_cleaned[encoded_column_name] = le.fit_transform(df_cleaned[col])

    # Store the encoder for possible future inverse transformation
    encoders[col] = le

    # Display actual value and encoded value side-by-side
    print(f"\nEncoding for column: {col}")
```

```

mapping = dict(zip(le.classes_, le.transform(le.classes_)))
for label, val in mapping.items():
    print(f" {label} -> {val}")

# Check the first few rows to confirm
#print("\nSample of DataFrame with original and encoded columns:")
#print(df_cleaned.head())
print("First 5 rows of the dataset:")
display(df_cleaned.head())

```

Encoding for column: gender

```

female -> 0
male -> 1
undefined -> 2

```

Encoding for column: stance

```

believer -> 0
denier -> 1
neutral -> 2

```

Encoding for column: topic

```

Donald Trump versus Science -> 0
Global stance -> 1
Ideological Positions on Global Warming -> 2
Impact of Resource Overconsumption -> 3
Importance of Human Intervantion -> 4
Politics -> 5
Seriousness of Gas Emissions -> 6
Significiance of Pollution Awareness Events -> 7
Undefined / One Word Hashtags -> 8
Weather Extremes -> 9

```

Encoding for column: aggressiveness

```

aggressive -> 0
not aggressive -> 1

```

First 5 rows of the dataset:

	id	created_at	lng	lat	sentiment	temperature_avg	gender	stance	topic	aggressiveness	gender_enc
1	13275	2006-07-23 21:52:30+00:00	-73.949582	40.650104	0.575777	-1.114768	undefined	neutral	Weather Extremes	aggressive	
7	1092823	2006-12-14 01:39:10+00:00	-122.419420	37.774930	-0.544195	4.228540	male	neutral	Ideological Positions on Global Warming	aggressive	
8	1278023	2006-12-17 19:43:09+00:00	-79.791980	36.072640	-0.565028	5.478175	male	denier	Weather Extremes	aggressive	
9	1455543	2006-12-21 01:39:01+00:00	-121.805790	38.004920	0.650960	-1.652156	male	neutral	Weather Extremes	not aggressive	
11	1893063	2006-12-31 10:47:25+00:00	-1.902691	52.479699	0.670905	4.864521	male	neutral	Weather Extremes	aggressive	



In [ ]:

In [6]: *#Scaling Numerical Features*

```
from sklearn.preprocessing import StandardScaler
```

```
scaler = StandardScaler()
```

```
df_cleaned[['sentiment', 'temperature_avg']] = scaler.fit_transform(df_cleaned[['sentiment', 'temperature_avg']])
```

In [7]: *#Dimensionality Reduction*

```
from sklearn.decomposition import PCA
```

```
# Selecting numeric columns for PCA
```

```
features = df_cleaned[['sentiment', 'temperature_avg']]
```

```
pca = PCA(n_components=2)
```

```
principal_components = pca.fit_transform(features)
```

```
df_cleaned[['PC1', 'PC2']] = principal_components
```



```
In [8]: #Data Wrangling Operations
# Clean category labels and lowercase
df_cleaned['gender'] = df_cleaned['gender'].astype(str).str.lower().str.strip()
df_cleaned['stance'] = df_cleaned['stance'].astype(str).str.lower().str.strip()
df_cleaned['topic'] = df_cleaned['topic'].astype(str).str.lower().str.strip()
```

```
In [ ]: # Preparation for Modeling
#X = df_cleaned.drop(columns=['aggressiveness']) # Features
#y = df_cleaned['aggressiveness']              # Target
```

```
In [11]: #Descriptive Analytics
# Shape of the dataset
print(f"Dataset contains {df_cleaned.shape[0]} rows and {df_cleaned.shape[1]} columns.")

# Data types and missing values
print("\nData info:")
df_cleaned.info()

# Summary statistics for numeric and categorical features
print("\nSummary statistics (numeric):")
display(df_cleaned.describe())

print("\nSummary statistics (categorical):")
display(df_cleaned.describe(include='object'))

# Check missing values
print("\nMissing values count:")
display(df_cleaned.isnull().sum())
```

Dataset contains 5307538 rows and 16 columns.

Data info:

```
<class 'pandas.core.frame.DataFrame'>
```

Index: 5307538 entries, 1 to 15789408

Data columns (total 16 columns):

#	Column	Dtype
0	id	int64
1	created_at	datetime64[ns, UTC]
2	lng	float64
3	lat	float64
4	sentiment	float64
5	temperature_avg	float64
6	gender	object
7	stance	object
8	topic	object
9	aggressiveness	object
10	gender_encoded	int64
11	stance_encoded	int64
12	topic_encoded	int64
13	aggressiveness_encoded	int64
14	PC1	float64
15	PC2	float64

dtypes: datetime64[ns, UTC](1), float64(6), int64(5), object(4)  
memory usage: 688.4+ MB

Summary statistics (numeric):

	id	lng	lat	sentiment	temperature_avg	gender_encoded	stance_encoded	topic_encoded	aggr
count	5.307538e+06	5.307538e+06	5.307538e+06	5.307538e+06	5.307538e+06	5.307538e+06	5.307538e+06	5.307538e+06	
mean	8.620935e+17	-4.639117e+01	3.408025e+01	-3.847011e-17	-4.896585e-17	7.179193e-01	4.437091e-01	4.198164e+00	
std	2.953852e+17	7.523162e+01	2.229430e+01	1.000000e+00	1.000000e+00	5.134762e-01	7.884864e-01	3.028200e+00	
min	1.327500e+04	-1.796670e+02	-9.000000e+01	-2.291631e+00	-6.456731e+00	0.000000e+00	0.000000e+00	0.000000e+00	
25%	7.805130e+17	-9.536327e+01	3.315067e+01	-9.138123e-01	-6.279653e-01	0.000000e+00	0.000000e+00	1.000000e+00	
50%	9.596823e+17	-7.703637e+01	3.995233e+01	-1.637279e-02	-8.851528e-03	1.000000e+00	0.000000e+00	4.000000e+00	
75%	1.049333e+18	-1.483154e-01	4.550884e+01	9.384739e-01	6.900383e-01	1.000000e+00	1.000000e+00	7.000000e+00	
max	1.178911e+18	1.793830e+02	8.500000e+01	2.206074e+00	5.199858e+00	2.000000e+00	2.000000e+00	9.000000e+00	



Summary statistics (categorical):

	gender	stance	topic	aggressiveness
count	5307538	5307538	5307538	5307538
unique	3	3	10	2
top	male	believer	global stance	not aggressive
freq	3485846	3947378	1462525	3774449

Missing values count:

	0
<b>id</b>	0
<b>created_at</b>	0
<b>lng</b>	0
<b>lat</b>	0
<b>sentiment</b>	0
<b>temperature_avg</b>	0
<b>gender</b>	0
<b>stance</b>	0
<b>topic</b>	0
<b>aggressiveness</b>	0
<b>gender_encoded</b>	0
<b>stance_encoded</b>	0
<b>topic_encoded</b>	0
<b>aggressiveness_encoded</b>	0
<b>PC1</b>	0
<b>PC2</b>	0

**dtype:** int64

```
In [ ]: # positive value means right
skewness = df_cleaned['sentiment'].skew()
print(f"Sentiment skewness: {skewness:.4f}")

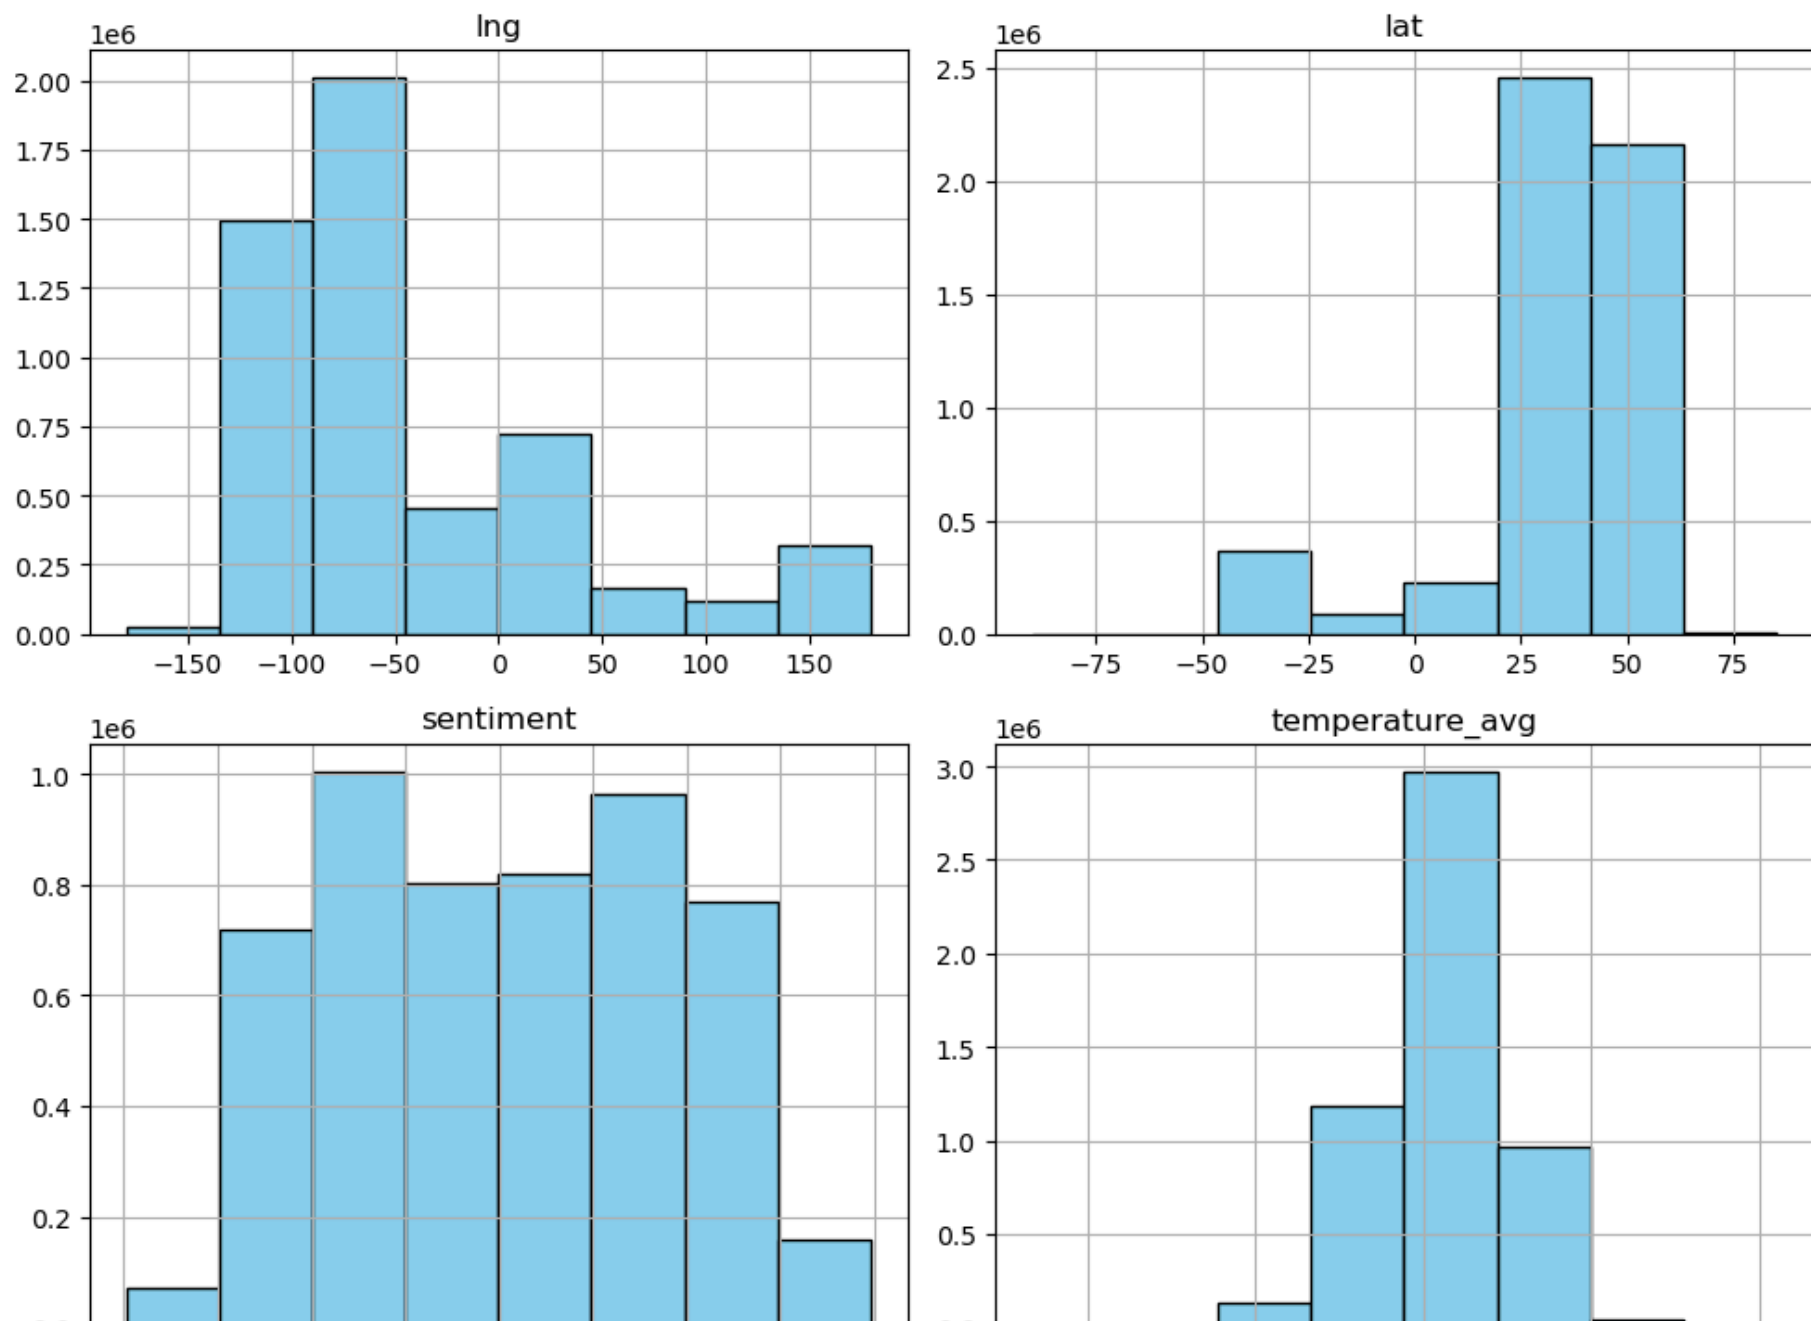
skewness_temp = df_cleaned['temperature_avg'].skew()
print(f"Temperature skewness: {skewness_temp:.4f}")
```

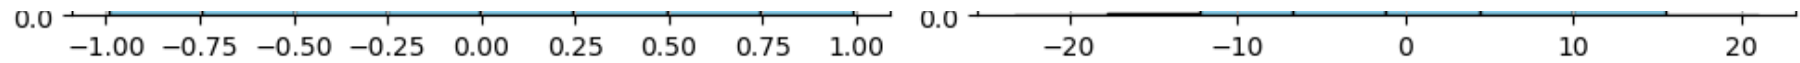
```
skewness_lng = df_cleaned['lng'].skew()  
print(f"Longitude skewness: {skewness_lng:.4f}")  
  
skewness_lat = df_cleaned['lat'].skew()  
print(f"Laitude skewness: {skewness_lat:.4f}")
```

Sentiment skewness: 0.0244  
Temparature skewness: -0.2713  
Longitude skewness: 1.2814  
Laitude skewness: -2.0892

```
In [ ]: #Histograms (Distribution Analysis)  
df_cleaned[['lng', 'lat', 'sentiment', 'temperature_avg']].hist(  
    bins=8, figsize=(10, 8), color='skyblue', edgecolor='black'  
)  
plt.suptitle("Histograms of Numerical Features", fontsize=16)  
plt.tight_layout()  
plt.show()
```

## Histograms of Numerical Features

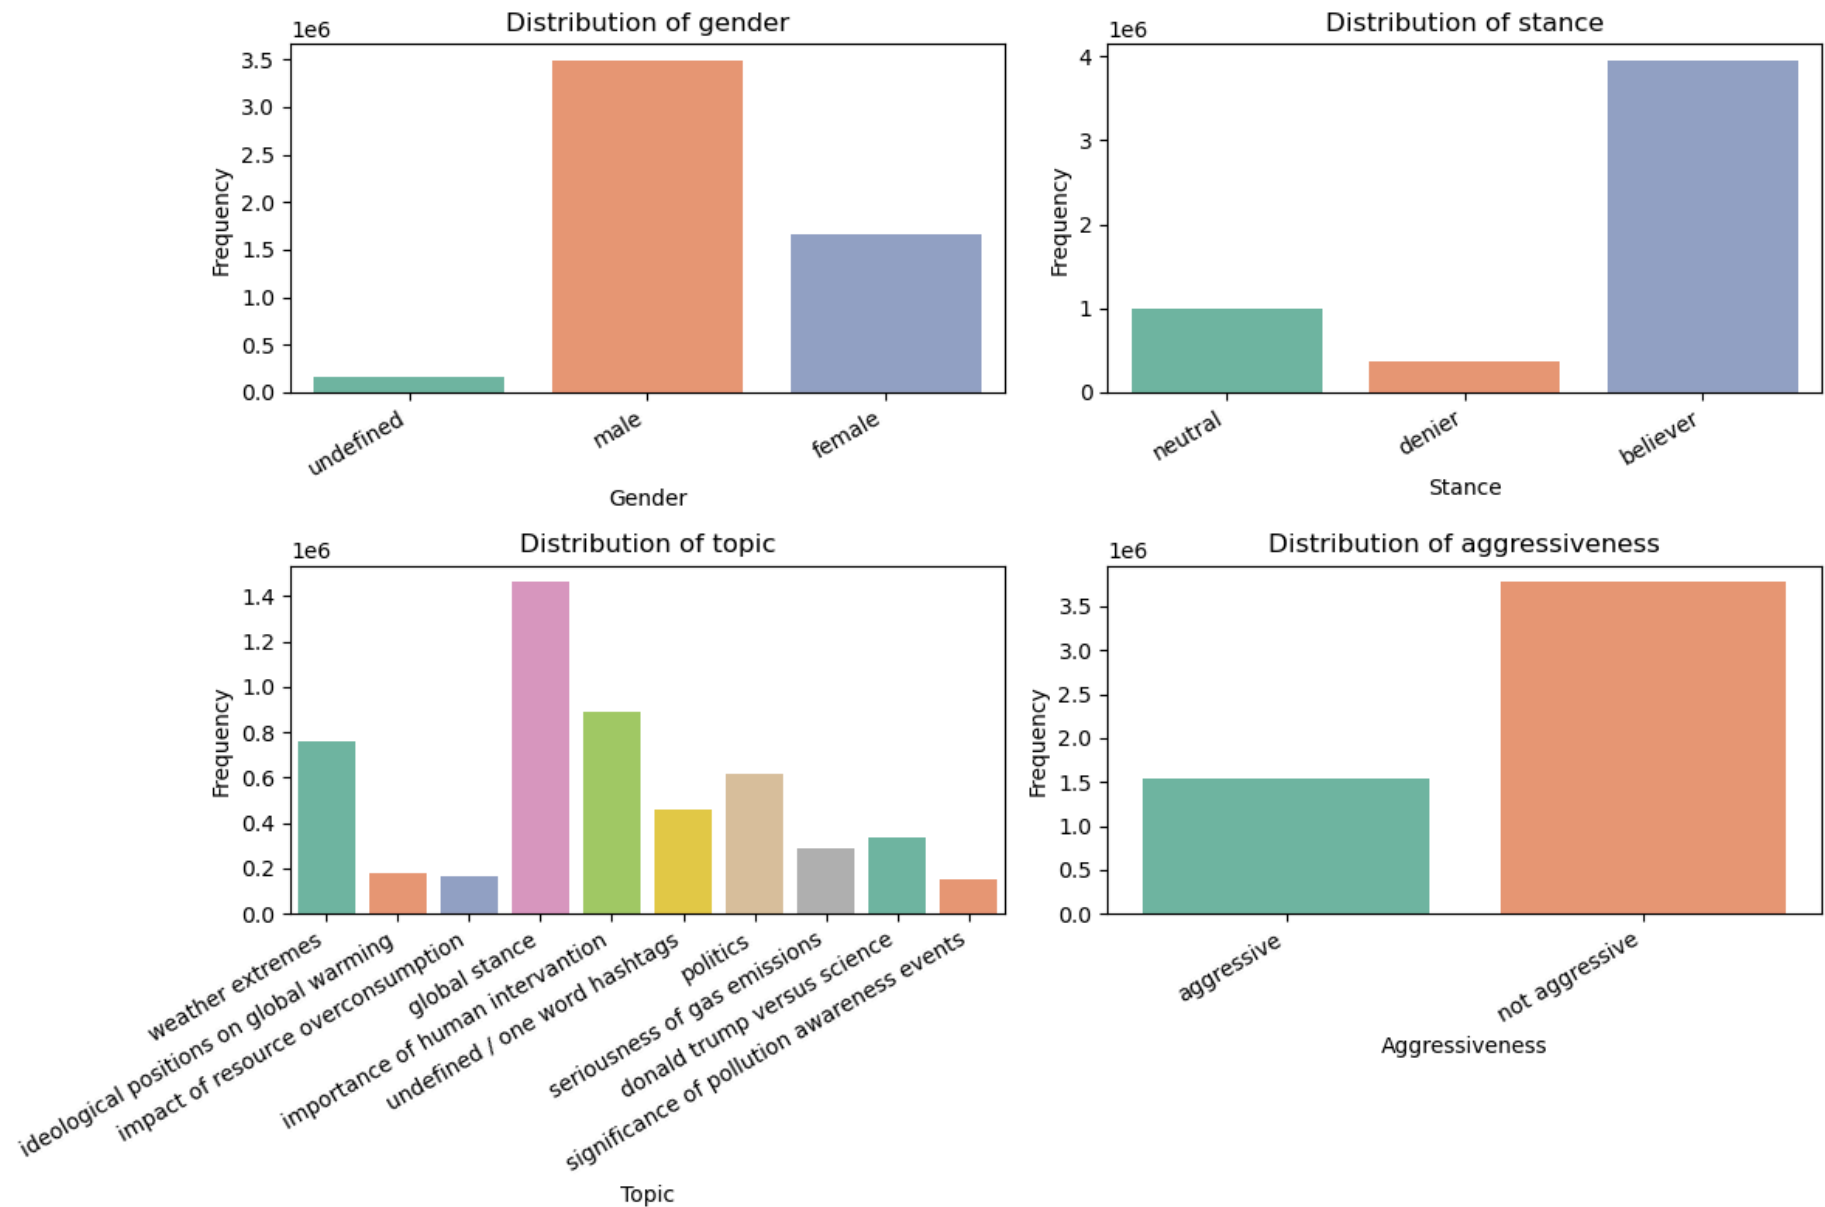




```
In [ ]: #bar chart Count of Categorical Variables
import seaborn as sns
import matplotlib.pyplot as plt

categorical_cols = ['gender', 'stance', 'topic', 'aggressiveness']

fig, axes = plt.subplots(2, 2, figsize=(12, 8))
for ax, col in zip(axes.flatten(), categorical_cols):
    sns.countplot(data=df_cleaned, x=col, hue=col, palette="Set2", legend=False, ax=ax)
    ax.set_title(f"Distribution of {col}")
    ax.set_ylabel("Frequency")
    ax.set_xlabel(col.capitalize())
    plt.setp(ax.get_xticklabels(), rotation=30, ha="right") # <- correct and safe way
plt.tight_layout()
plt.show()
```

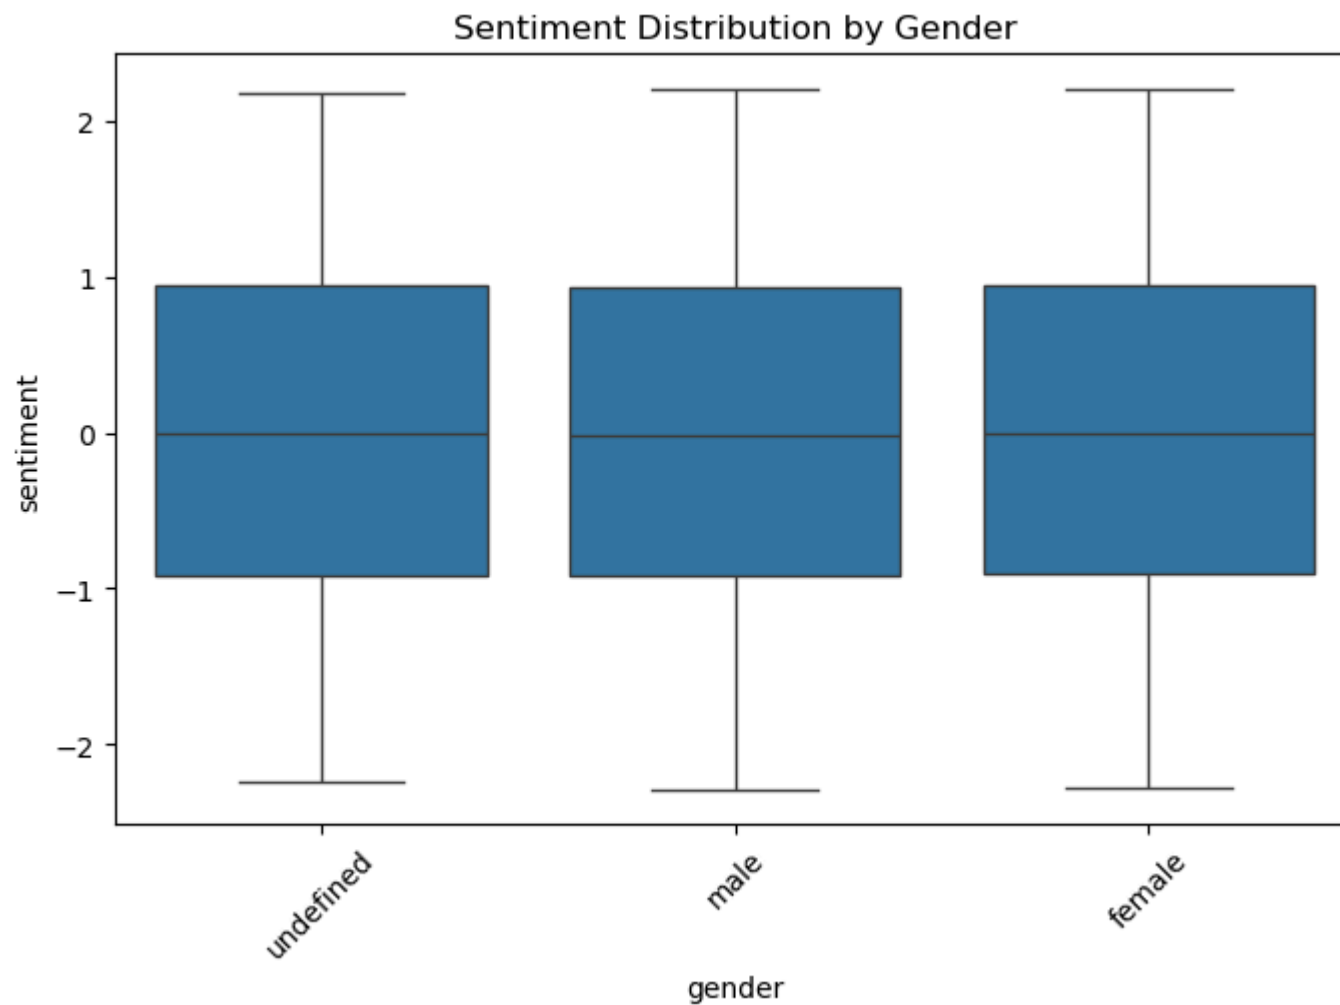


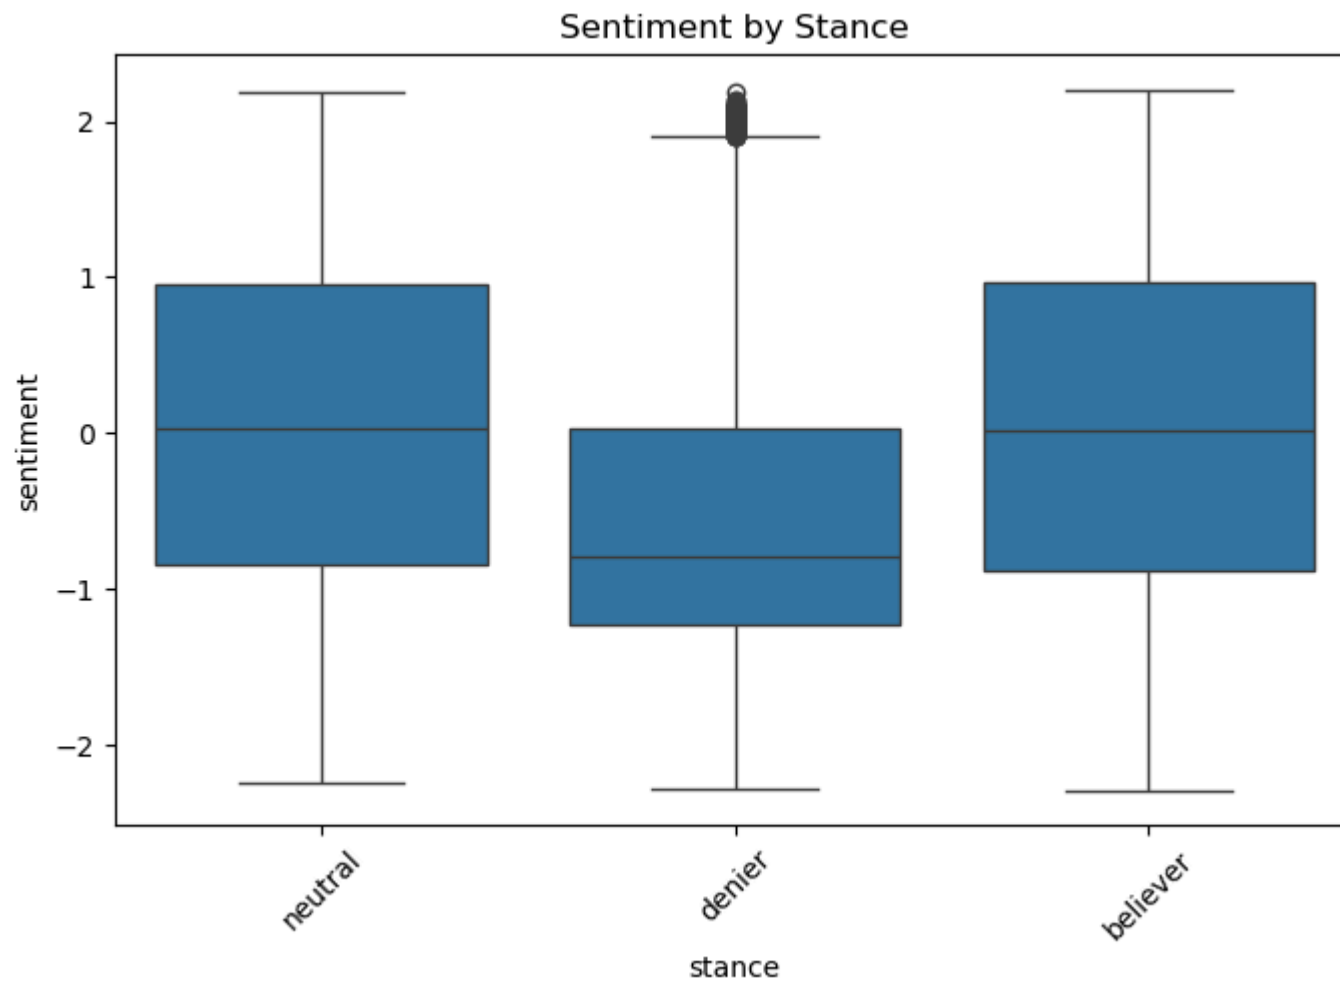
```
In [ ]: #Compare distributions across groups like gender and stance
plt.figure(figsize=(8, 5))
sns.boxplot(x='gender', y='sentiment', data=df_cleaned)
```



```
plt.title('Sentiment Distribution by Gender')
plt.xticks(rotation=45)
plt.show()

plt.figure(figsize=(8, 5))
sns.boxplot(x='stance', y='sentiment', data=df_cleaned)
plt.title('Sentiment by Stance')
plt.xticks(rotation=45)
plt.show()
```



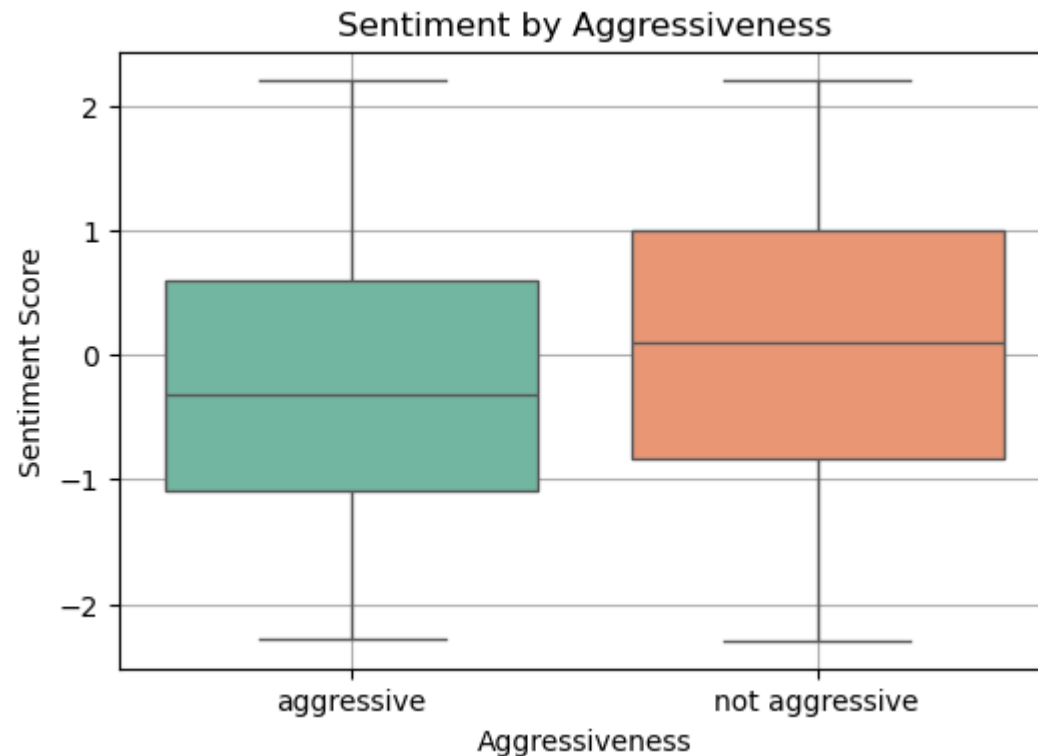


```
In [ ]: #Boxplot: Sentiment by Aggressiveness
plt.figure(figsize=(6, 4))
sns.boxplot(data=df_cleaned, x='aggressiveness', y='sentiment', palette='Set2')
plt.title('Sentiment by Aggressiveness')
plt.ylabel('Sentiment Score')
plt.xlabel('Aggressiveness')
plt.grid(True)
plt.show()
```

```
/tmp/ipykernel_8162/2581342232.py:3: FutureWarning:
```

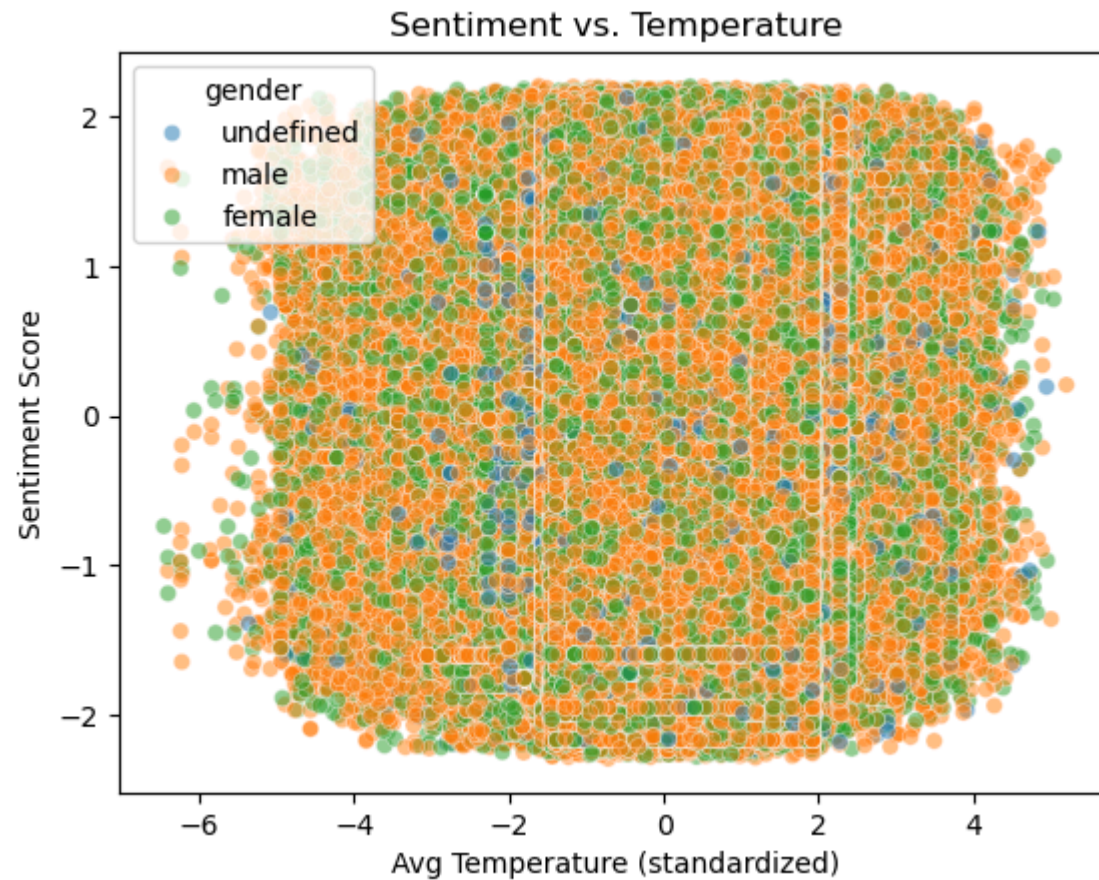
Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `x` variable to `hue` and set `legend=False` for the same effect.

```
sns.boxplot(data=df_cleaned, x='aggressiveness', y='sentiment', palette='Set2')
```



```
In [ ]: # Sentimenplt.figure(figsize=(8, 5))
sns.scatterplot(data=df_cleaned, x='temperature_avg', y='sentiment', hue='gender', alpha=0.5)
plt.title('Sentiment vs. Temperature')
plt.xlabel('Avg Temperature (standardized)')
plt.ylabel('Sentiment Score')
plt.show()
```

```
/home/opc/anaconda3/lib/python3.12/site-packages/IPython/core/pylabtools.py:170: UserWarning: Creating legend with loc="best" can be slow with large amounts of data.
fig.canvas.print_figure(bytes_io, **kw)
```

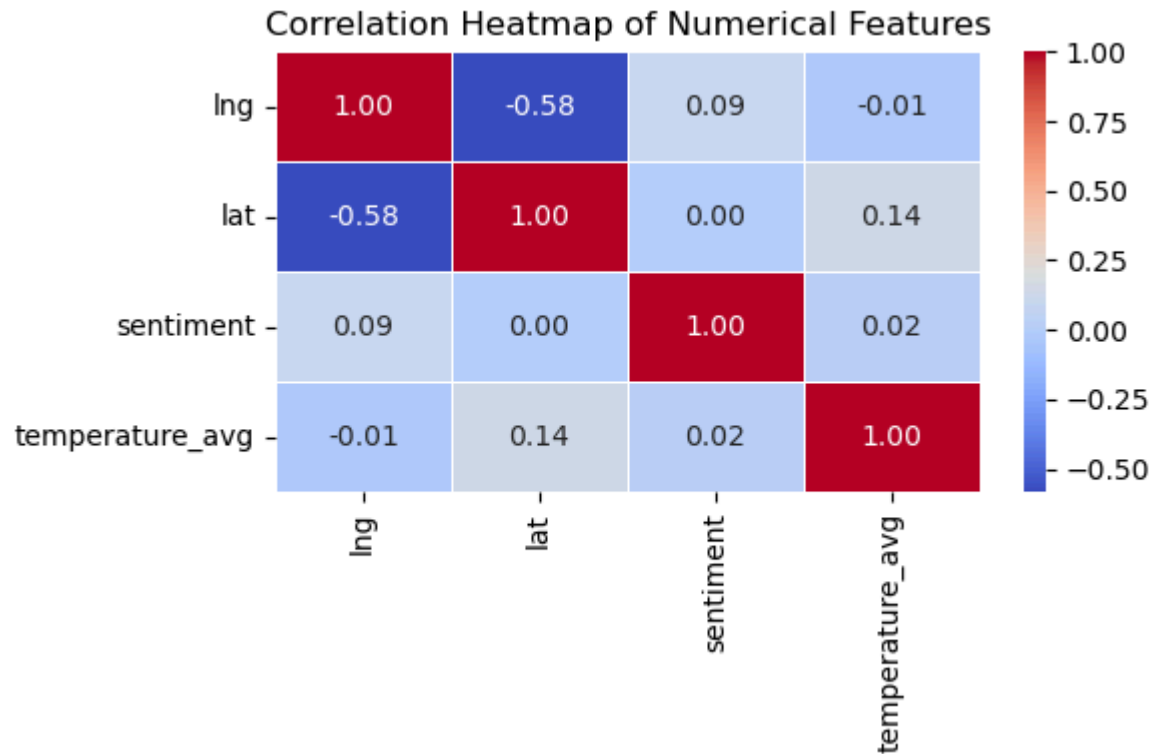


```
In [ ]: import seaborn as sns
import matplotlib.pyplot as plt

# Select only numerical columns
numeric_cols = ['lng', 'lat', 'sentiment', 'temperature_avg']
corr_matrix = df_cleaned[numeric_cols].corr()

# Create heatmap
plt.figure(figsize=(6, 4))
sns.heatmap(corr_matrix, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title('Correlation Heatmap of Numerical Features')
```

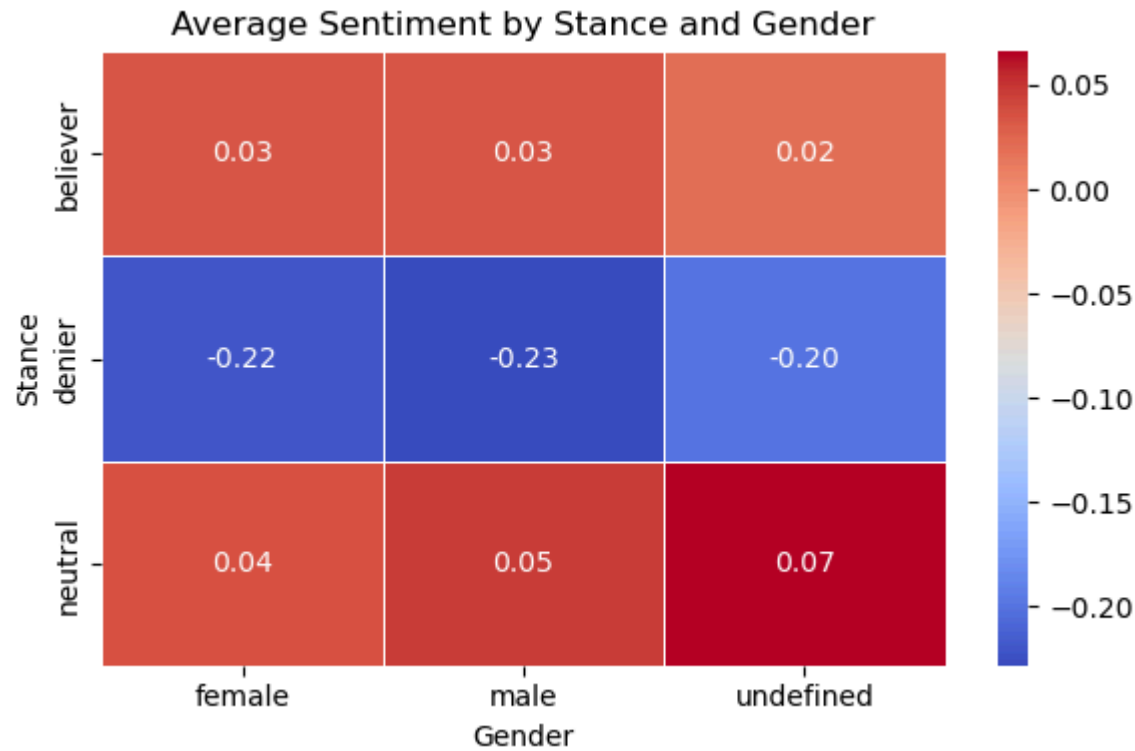
```
plt.tight_layout()
plt.show()
```



```
In [ ]: # Create pivot table from df_cleaned
pivot_table = df_cleaned.pivot_table(
    values='sentiment',
    index='stance',
    columns='gender',
    aggfunc='mean'
)

# Plot heatmap
plt.figure(figsize=(6, 4))
sns.heatmap(pivot_table, annot=True, cmap='coolwarm', fmt=".2f", linewidths=0.5)
plt.title("Average Sentiment by Stance and Gender")
plt.xlabel("Gender")
plt.ylabel("Stance")
```

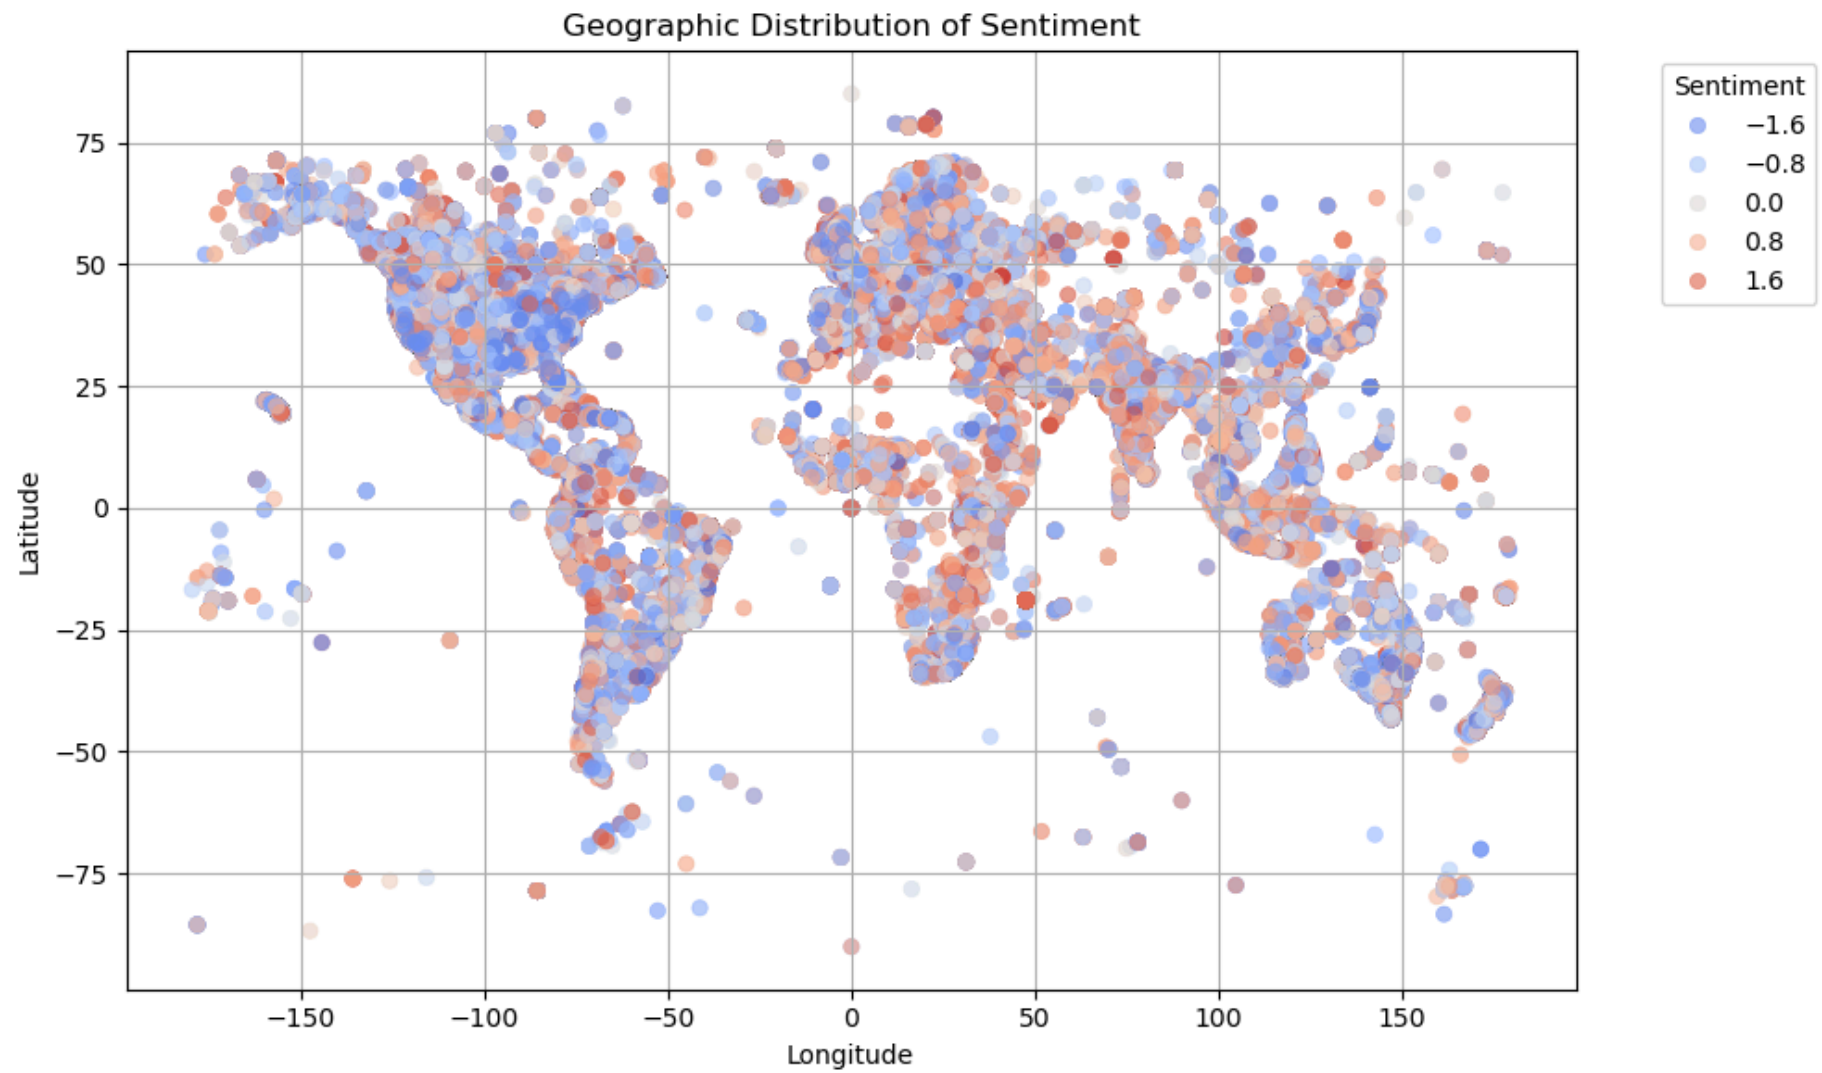
```
plt.tight_layout()
plt.show()
```



```
In [ ]: plt.figure(figsize=(10, 6))
sns.scatterplot(
    data=df_cleaned,
    x='lng',          # Longitude on X-axis
    y='lat',          # Latitude on Y-axis
    hue='sentiment',  # Color-coded by sentiment
    palette='coolwarm', # Blue to red color gradient
    alpha=0.6,
    edgecolor=None
)

plt.title('Geographic Distribution of Sentiment')
plt.xlabel('Longitude')
plt.ylabel('Latitude')
```

```
plt.legend(title='Sentiment', bbox_to_anchor=(1.05, 1), loc='upper left')  
plt.grid(True)  
plt.tight_layout()  
plt.show()
```



```
In [ ]: #descriptive for categorical feautres - Frequency Counts (Distribution)  
categorical_cols = ['gender', 'stance', 'topic', 'aggressiveness']  
for col in categorical_cols:
```

```
print(f"\n{col} value counts:\n")
display(df_cleaned[col].value_counts())
```

gender value counts:

```
gender
male      3485846
female    1659423
undefined  162269
Name: count, dtype: int64
stance value counts:
```

```
stance
believer  3947378
neutral   994843
denier     365317
Name: count, dtype: int64
topic value counts:
```

```
topic
global stance      1462525
importance of human intervantion  889463
weather extremes   761129
politics           618945
undefined / one word hashtags    458904
donald trump versus science      333777
seriousness of gas emissions     291323
ideological positions on global warming  176568
impact of resource overconsumption  164996
significance of pollution awareness events  149908
Name: count, dtype: int64
aggressiveness value counts:
```

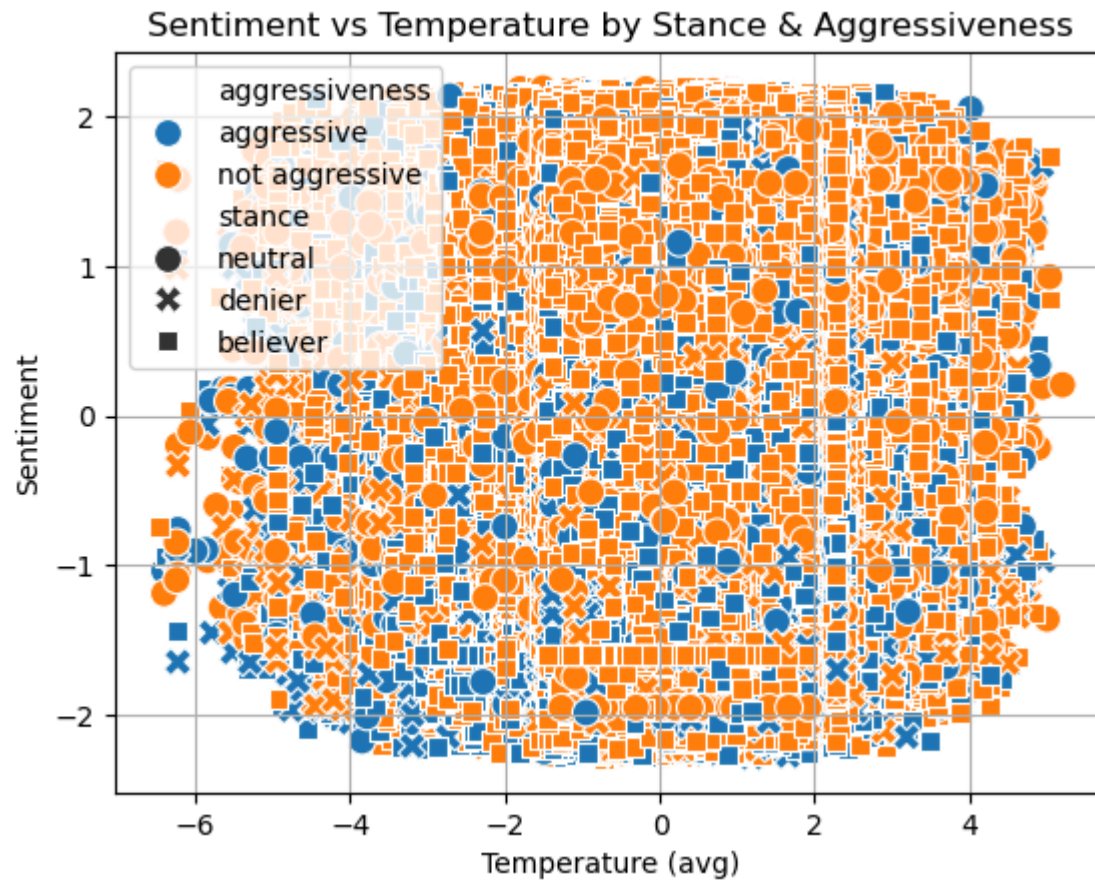
```
aggressiveness
not aggressive  3774449
aggressive      1533089
Name: count, dtype: int64
```

```
In [ ]: #Sentiment vs. Temperature (with Aggressiveness)
sns.scatterplot(
    data=df_cleaned,
```



```
x='temperature_avg',  
y='sentiment',  
hue='aggressiveness',  
style='stance',  
s=100  
)  
plt.title("Sentiment vs Temperature by Stance & Aggressiveness")  
plt.xlabel("Temperature (avg)")  
plt.ylabel("Sentiment")  
plt.grid(True)  
plt.show()
```

```
/home/opc/anaconda3/lib/python3.12/site-packages/IPython/core/pylabtools.py:170: UserWarning: Creating legend with loc="best" can be slow with large amounts of data.  
fig.canvas.print_figure(bytes_io, **kw)
```

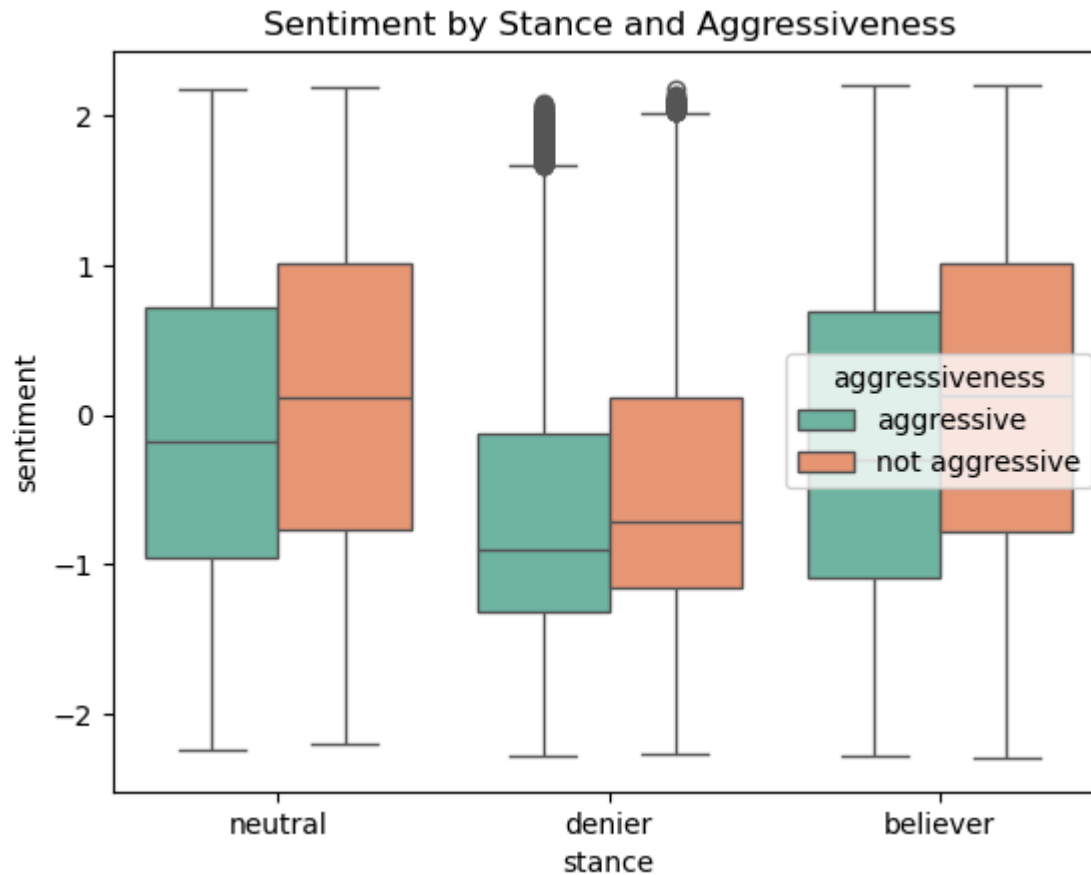


```
In [ ]: #Cross-tabulations (e.g., stance vs aggressiveness)
pd.crosstab(df_cleaned['stance'], df_cleaned['aggressiveness'], normalize='index') * 100
```

```
Out[ ]: aggressiveness  aggressive  not aggressive
```

stance		
believer	27.497671	72.502329
denier	42.678277	57.321723
neutral	29.325331	70.674669

```
In [ ]: #Boxplots for Numerical Variables by Category
sns.boxplot(data=df_cleaned, x='stance', y='sentiment', hue='aggressiveness', palette='Set2')
plt.title("Sentiment by Stance and Aggressiveness")
plt.show()
```



```
In [ ]: #Diagnostic Analytics - Hypothesis Testing - T-Test- Does gender affect sentiment?
#If  $p < 0.05$ , then gender has a significant impact on sentiment.
from scipy.stats import ttest_ind

# Extract sentiment scores by gender
male = df_cleaned[df_cleaned['gender'] == 'male']['sentiment'].dropna()
female = df_cleaned[df_cleaned['gender'] == 'female']['sentiment'].dropna()
```

```

print("Male count:", len(male))
print("Female count:", len(female))

# Perform t-test assuming unequal variances (t-test)
t_stat, p_val = ttest_ind(male, female, equal_var=False)

print(f"T-Test: t={t_stat:.4f}, p={p_val:.4f}")

```

Male count: 3485846

Female count: 1659423

T-Test: t=-11.5635, p=0.0000

```

In [ ]: import matplotlib.pyplot as plt
import seaborn as sns
import pandas as pd

# Convert created_at to datetime
df['created_at'] = pd.to_datetime(df_cleaned['created_at'], errors='coerce')
df = df.dropna(subset=['created_at', 'stance'])

# Create a new column for year-month
df['year_month'] = df['created_at'].dt.to_period('M')
#df['year_month'] = df['created_at'].dt.strftime('%Y-%b').str.upper()

# Group by month and stance
stance_trends = df.groupby(['year_month', 'stance']).size().unstack(fill_value=0)

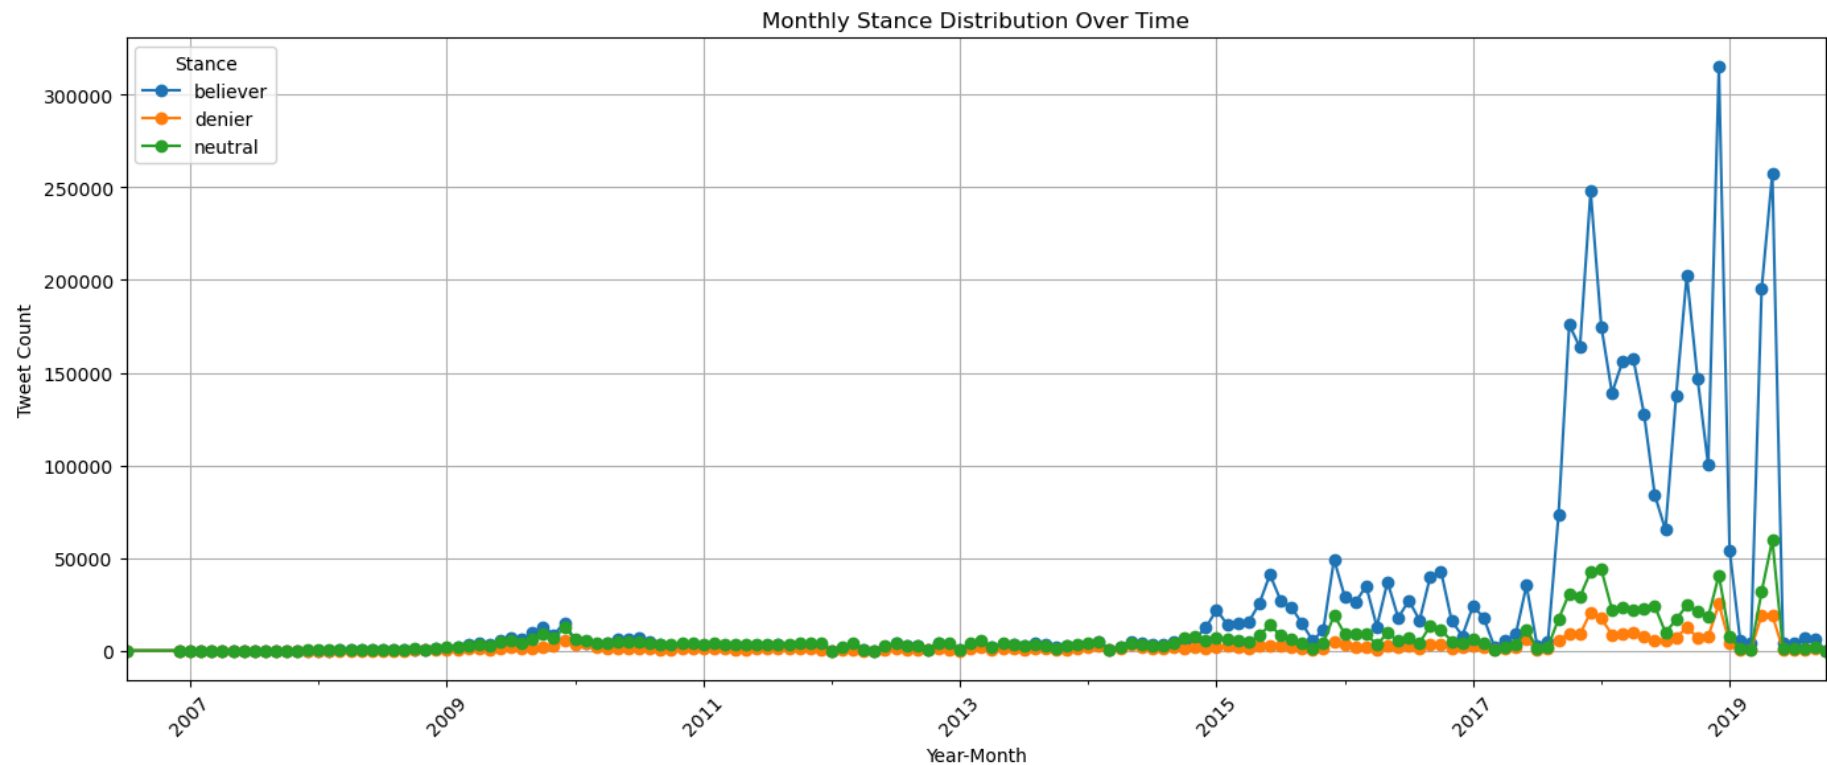
# Plot stance trends over time
plt.figure(figsize=(14, 6))
stance_trends.plot(marker='o', figsize=(14, 6))
plt.title('Monthly Stance Distribution Over Time')
plt.xlabel('Year-Month')
plt.ylabel('Tweet Count')
plt.xticks(rotation=45)
plt.grid(True)
plt.tight_layout()
plt.legend(title='Stance')
plt.show()

```

```
/tmp/ipykernel_299505/1667700143.py:10: UserWarning: Converting to PeriodArray/Index representation will drop timezone information.
```

```
df['year_month'] = df['created_at'].dt.to_period('M')
```

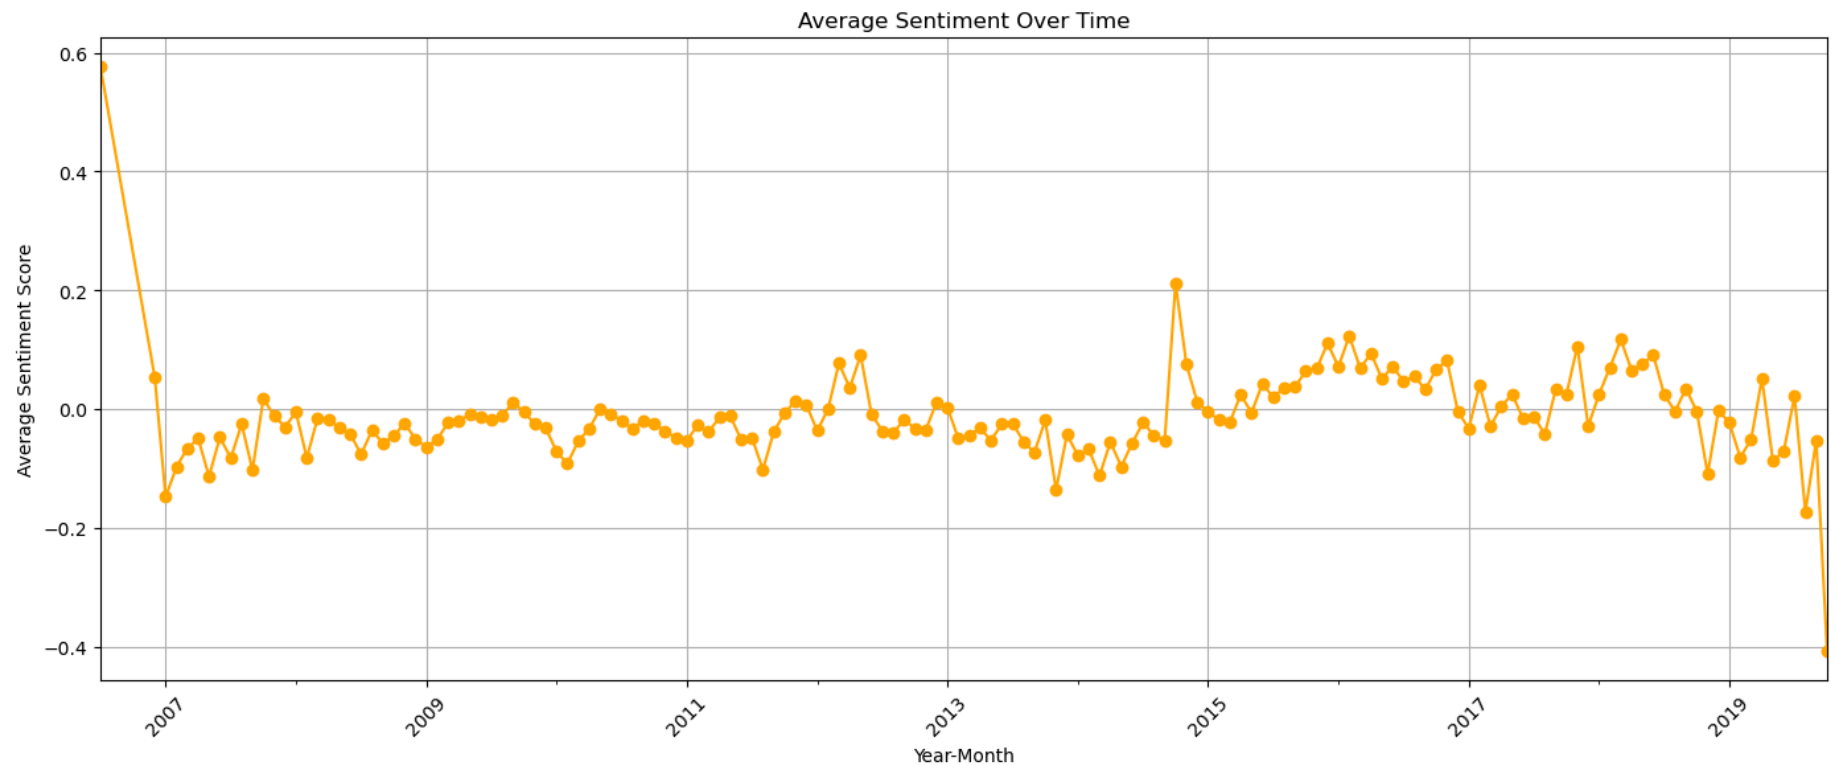
```
<Figure size 1400x600 with 0 Axes>
```



```
In [ ]: # Group by year-month and calculate average sentiment
sentiment_trends = df.groupby('year_month')['sentiment'].mean()

# Plot sentiment trends over time
plt.figure(figsize=(14, 6))
sentiment_trends.plot(marker='o', color='orange')
plt.title('Average Sentiment Over Time')
plt.xlabel('Year-Month')
plt.ylabel('Average Sentiment Score')
plt.xticks(rotation=45)
plt.grid(True)
```

```
plt.tight_layout()
plt.show()
```



```
In [ ]: # NOVA - Sentiment across stances
# A significant p-value suggests sentiment varies by stance
from scipy.stats import f_oneway

groups = [group['sentiment'].dropna() for name, group in df_cleaned.groupby('stance')]
f_stat, p_val = f_oneway(*groups)
print(f"ANOVA: F={f_stat:.4f}, p={p_val:.4f}")
```

ANOVA: F=61007.3680, p=0.0000

```
In [ ]: #Diagnostic Regression - Sentiment as Dependent Variable
#R² tells how much variance is explained by predictors.
#Feature weights can suggest direction of influence (positive/negative).
from sklearn.linear_model import LinearRegression
```

```
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import LabelEncoder, StandardScaler
from sklearn.metrics import r2_score, root_mean_squared_error
import pandas as pd

# Prepare the data
diag_df = df_cleaned[['sentiment', 'temperature_avg', 'gender_encoded', 'stance_encoded', 'topic_encoded']].dropna()

# Encode categorical features
#for col in ['gender', 'stance', 'topic']:
#    diag_df[col] = LabelEncoder().fit_transform(diag_df[col])

# Define features and target
X = diag_df.drop(columns='sentiment')
y = diag_df['sentiment']

# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# Split data into train and test sets
X_train, X_test, y_train, y_test = train_test_split(
    X_scaled, y, test_size=0.3, random_state=42
)

# Train linear regression model
model = LinearRegression()
model.fit(X_train, y_train)

# Predict on test set
y_pred = model.predict(X_test)

# Evaluate the model
r2 = r2_score(y_test, y_pred)
rmse = root_mean_squared_error(y_test, y_pred)

print(f"R²: {r2:.4f}")
print(f"RMSE: {rmse:.4f}")
```

R²: 0.0008

RMSE: 0.4406

```
In [ ]: # Diagnostic Visualization & Feature Importance for Stance Prediction

import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.preprocessing import StandardScaler

# 1. Define features (exclude 'stance_encoded') and target = 'stance_encoded'
features = [
    'sentiment',
    'temperature_avg',
    'gender_encoded',
    'aggressiveness_encoded', # now used as a predictor
    'topic_encoded'
]
X_diag = df_cleaned[features].dropna()
y_diag = df_cleaned.loc[X_diag.index, 'stance_encoded']

# 2. Standardize numeric features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X_diag)
X_diag = pd.DataFrame(X_scaled, columns=features)

# 3. Train/test split
X_train, X_test, y_train, y_test = train_test_split(
    X_diag, y_diag, test_size=0.3, random_state=42, stratify=y_diag
)

# 4. Train Random Forest
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train, y_train)

# 5. Extract feature importances
importances = rf_model.feature_importances_
feature_names = features

# 6. Map to readable labels
labels = {
```

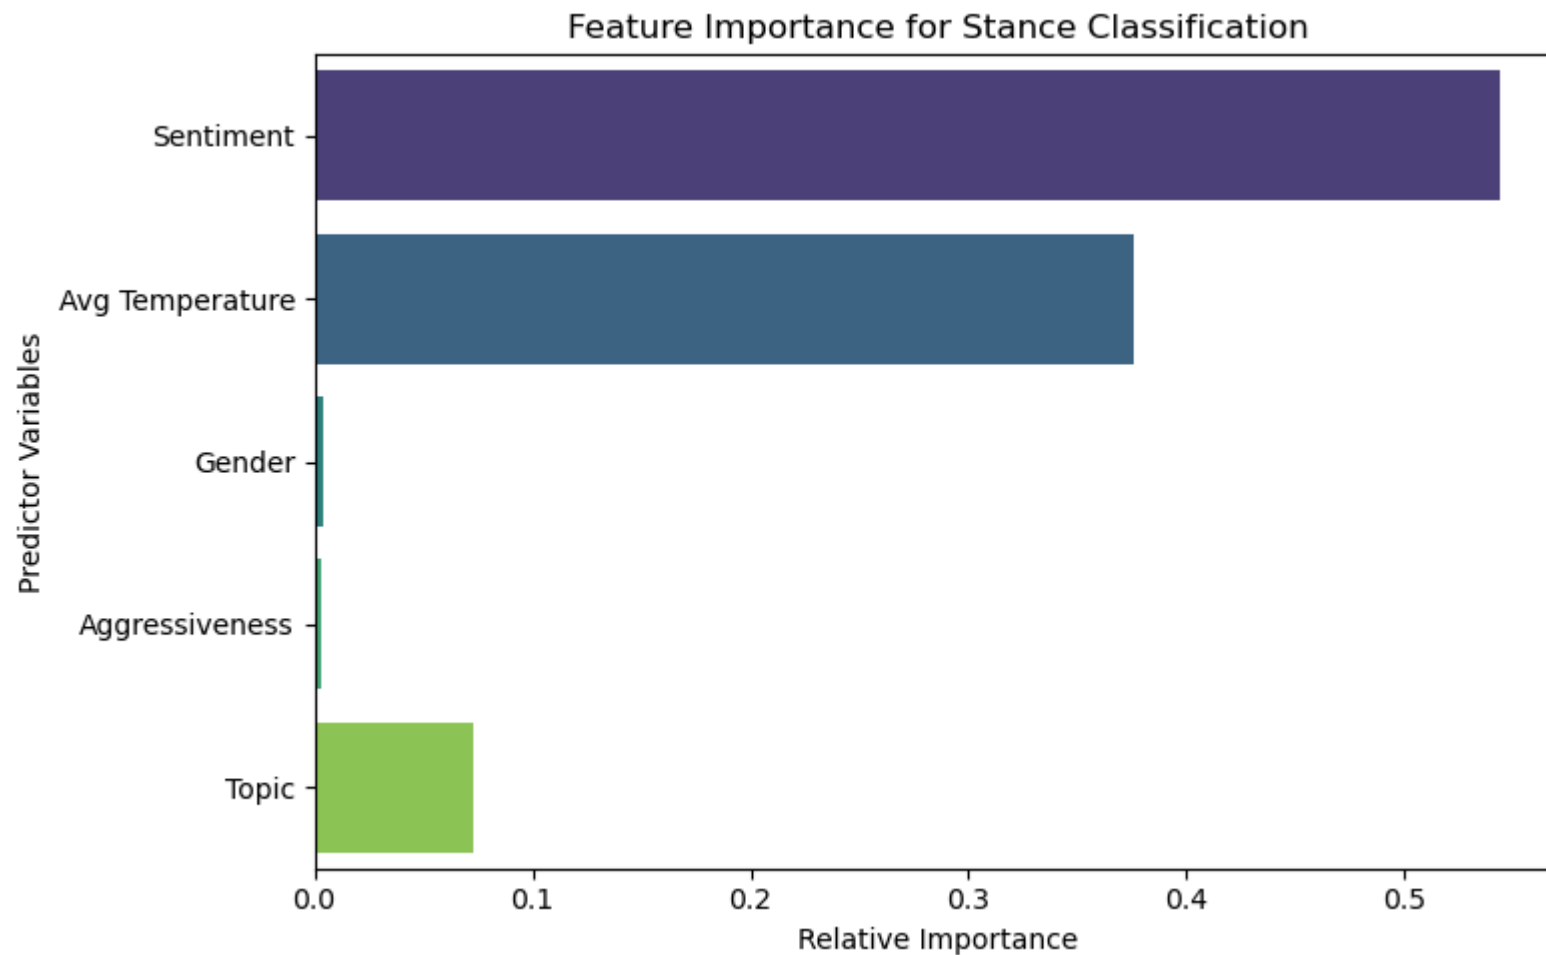


```
'sentiment': 'Sentiment',  
'temperature_avg': 'Avg Temperature',  
'gender_encoded': 'Gender',  
'aggressiveness_encoded': 'Aggressiveness',  
'topic_encoded': 'Topic'  
}  
readable = [labels[f] for f in feature_names]  
  
# 7. Plot  
plt.figure(figsize=(8, 5))  
sns.barplot(x=importances, y=readable, palette="viridis")  
plt.title("Feature Importance for Stance Classification")  
plt.xlabel("Relative Importance")  
plt.ylabel("Predictor Variables")  
plt.tight_layout()  
plt.show()
```

/tmp/ipykernel\_52256/2863863304.py:51: FutureWarning:

Passing `palette` without assigning `hue` is deprecated and will be removed in v0.14.0. Assign the `y` variable to `hue` and set `legend=False` for the same effect.

```
sns.barplot(x=importances, y=readable, palette="viridis")
```



```
In [ ]: from sklearn.tree import DecisionTreeClassifier, plot_tree
        from sklearn.preprocessing import LabelEncoder
        import matplotlib.pyplot as plt

        # 1. Encode stance if not already encoded
        #le_stance = LabelEncoder()
        #df_cleaned['stance_encoded'] = le_stance.fit_transform(df_cleaned['stance'])
        #tance_classes = list(le_stance.classes_) # e.g. ['believer', 'denier', 'neutral']

        # 2. Define features (exclude 'stance_encoded') and new target
        features = ['sentiment', 'temperature_avg', 'gender_encoded', 'aggressiveness_encoded', 'topic_encoded']
```

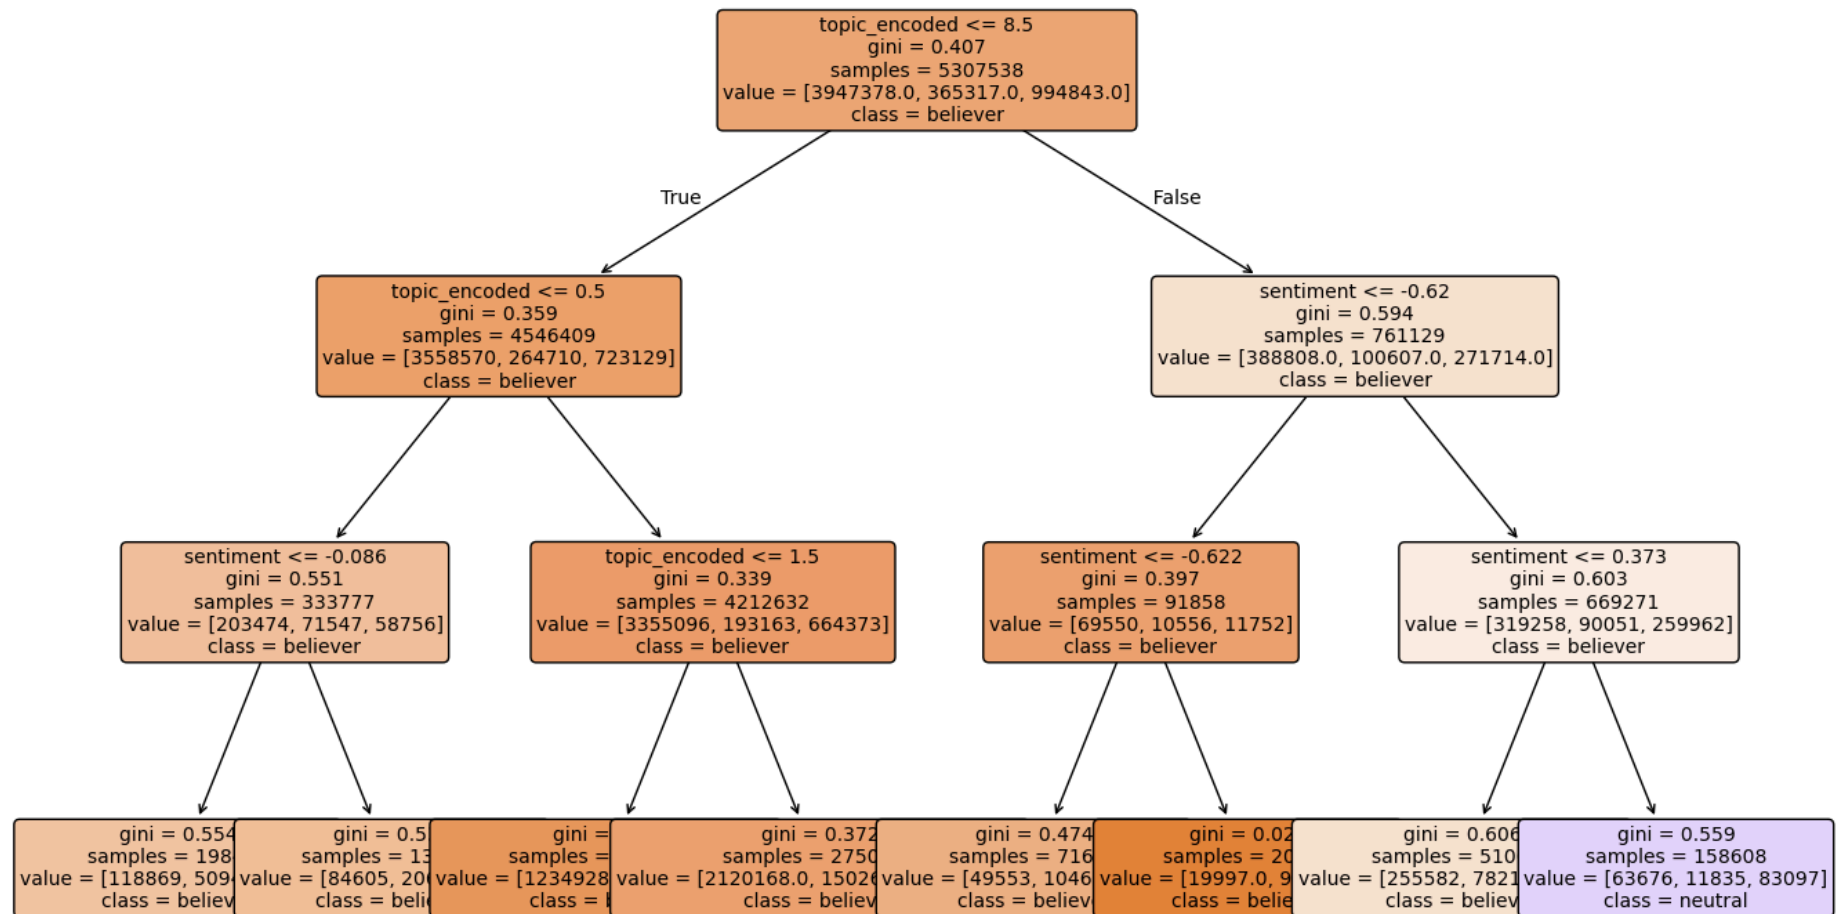
```
X_cls = df_cleaned[features].dropna()
y_cls = df_cleaned.loc[X_cls.index, 'stance_encoded']

# 3. Train the Decision Tree
tree_model = DecisionTreeClassifier(max_depth=3, random_state=42)
tree_model.fit(X_cls, y_cls)

# 4. Plot the tree
plt.figure(figsize=(16, 10))
plot_tree(
    tree_model,
    feature_names=features,
    class_names= ['believer', 'denier', 'neutral'], # stance_classes,
    filled=True,
    rounded=True,
    fontsize=10
)
plt.title("Decision Tree (Depth=3) 🌳 Stance Classification")
plt.show()
```

```
/home/opc/anaconda3/lib/python3.12/site-packages/IPython/core/pylabtools.py:170: UserWarning: Glyph 150 (\x96) missing from font(s) DejaVu Sans.
  fig.canvas.print_figure(bytes_io, **kw)
```

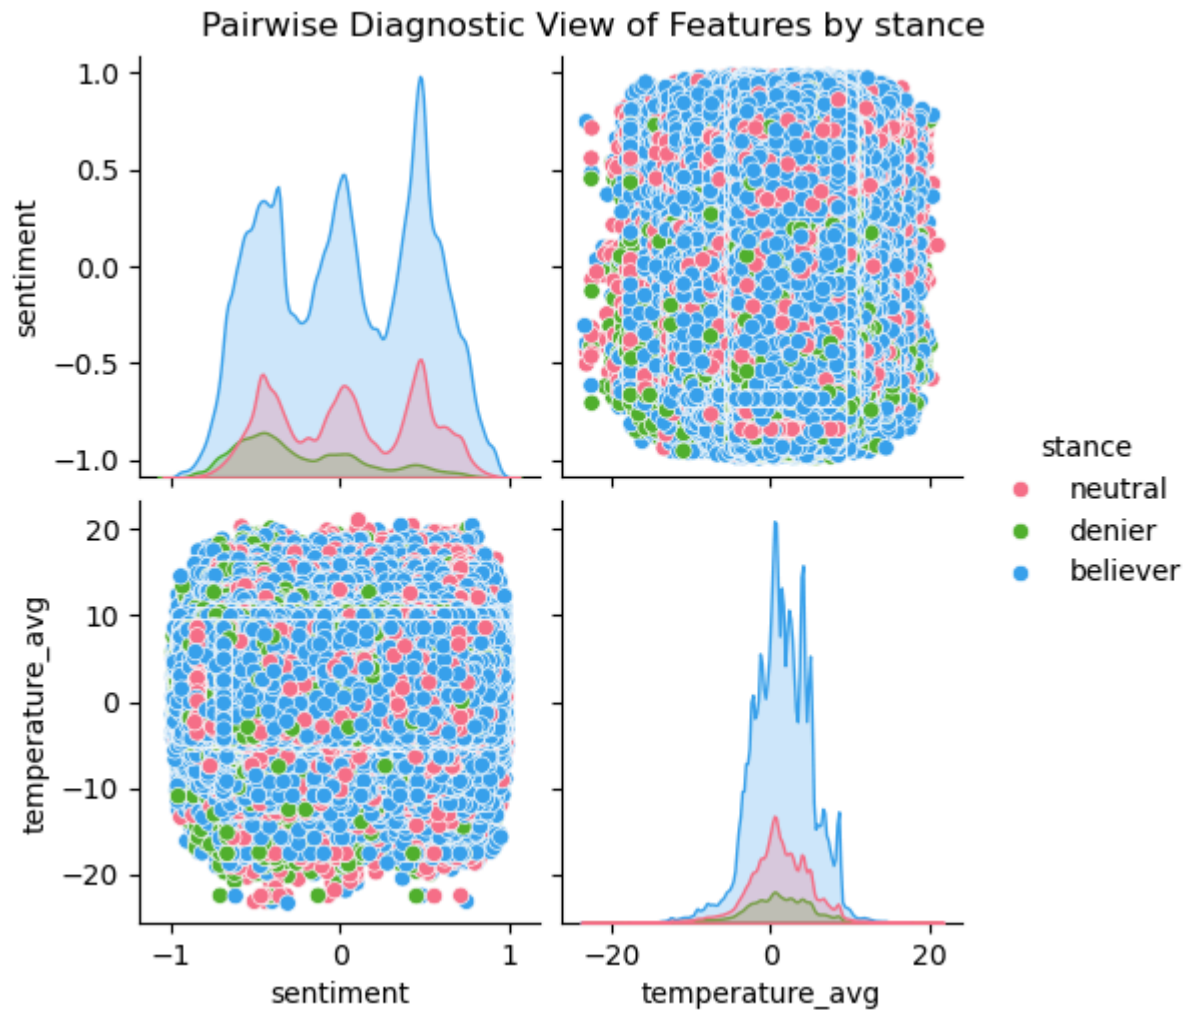
Decision Tree (Depth=3) □ Stance Classification



```

In [ ]: #Pairplot with Hue by stance - Visualize pairwise feature relationships, colored by target class.
sns.pairplot(df_cleaned, vars=['sentiment', 'temperature_avg'], hue='stance', palette='husl')
plt.suptitle('Pairwise Diagnostic View of Features by stance', y=1.02)
plt.show()

```



```
In [ ]: # Predictive Analytics
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
#from sklearn.metrics import classification_report, confusion_matrix
from sklearn.preprocessing import StandardScaler

from sklearn.metrics import (
    classification_report,
    confusion_matrix,
```

```
ConfusionMatrixDisplay,  
roc_curve,  
roc_auc_score  
)  
import matplotlib.pyplot as plt  
import seaborn as sns  
from sklearn.preprocessing import label_binarize  
from sklearn.metrics import roc_curve, auc  
  
# Load the full sheet  
#df = xls.parse("Sheet1")  
  
# Define features and target  
features = ['sentiment', 'temperature_avg', 'gender_encoded', 'topic_encoded', 'aggressiveness_encoded', 'PC1', 'PC2']  
target = 'stance_encoded'  
  
# Drop rows with missing target or feature values  
df_model = df_cleaned.dropna(subset=features + [target])  
  
# Split into train and test sets  
X = df_model[features]  
y = df_model[target]  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42)  
  
# Standardize features  
scaler = StandardScaler()  
X_train_scaled = scaler.fit_transform(X_train)  
X_test_scaled = scaler.transform(X_test)  
  
# Train a Random Forest Classifier  
clf = RandomForestClassifier(random_state=42)  
clf.fit(X_train_scaled, y_train)  
  
# Make predictions  
y_pred = clf.predict(X_test_scaled)  
  
# Evaluate the model  
#report = classification_report(y_test, y_pred, output_dict=True)  
#conf_matrix = confusion_matrix(y_test, y_pred)  
  
#report, conf_matrix
```

```

# Classification Report
stance_labels = ['believer', 'neutral', 'denier']
print("Classification Report:\n")
print(classification_report(y_test, y_pred, target_names=stance_labels))

# Confusion Matrix
cm = confusion_matrix(y_test, y_pred)
disp = ConfusionMatrixDisplay(confusion_matrix=cm, display_labels=stance_labels)

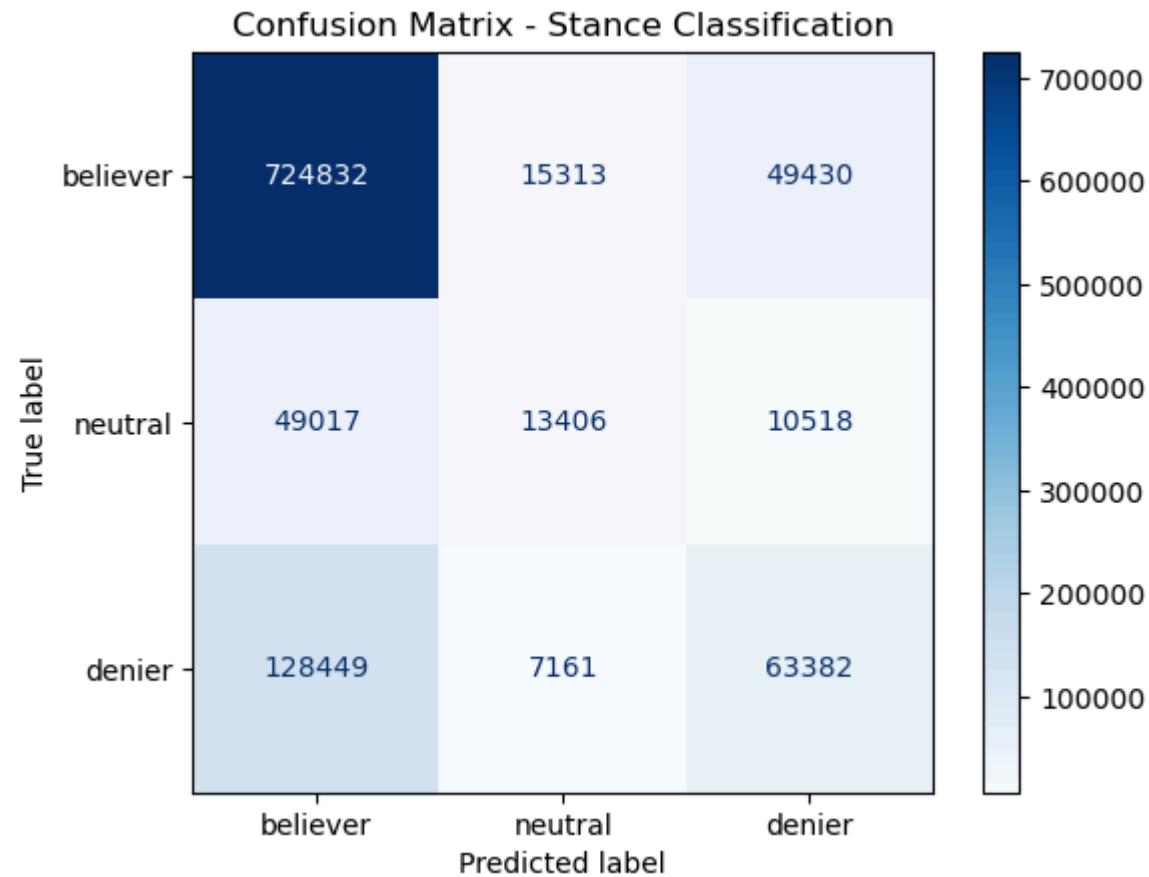
plt.figure(figsize=(6, 5))
disp.plot(cmap="Blues", values_format='d')
plt.title("Confusion Matrix - Stance Classification")
plt.grid(False)
plt.show()

```

Classification Report:

	precision	recall	f1-score	support
believer	0.80	0.92	0.86	789575
neutral	0.37	0.18	0.25	72941
denier	0.51	0.32	0.39	198992
accuracy			0.76	1061508
macro avg	0.56	0.47	0.50	1061508
weighted avg	0.72	0.76	0.73	1061508

<Figure size 600x500 with 0 Axes>



```
In [13]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.ensemble import RandomForestClassifier
from sklearn.metrics import classification_report, confusion_matrix, roc_curve, auc
from sklearn.preprocessing import StandardScaler, label_binarize
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.over_sampling import SMOTE

# Features and target
features = ['sentiment', 'temperature_avg', 'gender_encoded', 'aggressiveness_encoded']
target = 'stance_encoded'
```



```
X = df_cleaned[features]
y = df_cleaned[target]

# Standardize features
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)

# SMOTE oversampling
smote = SMOTE(random_state=42)
X_resampled, y_resampled = smote.fit_resample(X_scaled, y)

# Train-test split
X_train, X_test, y_train, y_test = train_test_split(
    X_resampled, y_resampled, test_size=0.3, random_state=42, stratify=y_resampled
)

# Train classifier --estimator set to 100, fist tried with diff value due to low cpu/gpu power
clf = RandomForestClassifier(n_estimators=100, max_depth=None, n_jobs=2, class_weight='balanced', random_state=42)
clf.fit(X_train, y_train)

# Predictions
y_pred = clf.predict(X_test)
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
print("\nClassification Report:\n", classification_report(y_test, y_pred))

# Confusion Matrix Plot
labels = ['Believer', 'Neutral', 'Denier']
cm = confusion_matrix(y_test, y_pred)
plt.figure(figsize=(6, 5))
sns.heatmap(cm, annot=True, fmt='d', cmap='Blues', xticklabels=labels, yticklabels=labels)
plt.xlabel('Predicted')
plt.ylabel('Actual')
plt.title('Confusion Matrix - Stance Classification')
plt.tight_layout()
plt.show()

# ROC Curve and AUC
y_test_bin = label_binarize(y_test, classes=np.unique(y))
y_score = clf.predict_proba(X_test)
n_classes = y_score.shape[1]
```

```

fpr, tpr, roc_auc = {}, {}, {}
for i in range(n_classes):
    fpr[i], tpr[i], _ = roc_curve(y_test_bin[:, i], y_score[:, i])
    roc_auc[i] = auc(fpr[i], tpr[i])

plt.figure(figsize=(8, 6))
colors = ['blue', 'green', 'red']
for i, color in zip(range(n_classes), colors):
    plt.plot(fpr[i], tpr[i], color=color, lw=2,
             label=f'{labels[i]} (AUC = {roc_auc[i]:.2f})')

plt.plot([0, 1], [0, 1], 'k--')
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve - Stance Classification')
plt.legend(loc="lower right")
plt.grid()
plt.tight_layout()
plt.show()

```

Confusion Matrix:

```

[[739691 194657 249865]
 [140555 930656 113003]
 [211841 139209 833164]]

```

Classification Report:

	precision	recall	f1-score	support
0	0.68	0.62	0.65	1184213
1	0.74	0.79	0.76	1184214
2	0.70	0.70	0.70	1184214
accuracy			0.70	3552641
macro avg	0.70	0.70	0.70	3552641
weighted avg	0.70	0.70	0.70	3552641

